

The Discovery Engine: Automated Theorem and Invariant Discovery in Judgment Geometry

Halley Young
Microsoft Research
halleyyoung@microsoft.com

March 22, 2026

Abstract

As software systems grow in scale, the bottleneck in formal verification shifts from proof construction to *invariant discovery*: finding the right properties to state before attempting to prove them. We present the JUGE Discovery Engine, a system that automatically discovers useful invariants and theorems about programs by exploring the semantic site. The engine orchestrates three cooperating subsystems: *novelty search* (selecting surprising properties not already in the portfolio), *analogy transport* (transferring proven properties from structurally similar programs), and *kind discovery* (classifying found properties by semantic category). A federation layer distributes discovery across large codebases. We prove a *completeness theorem*: for any finite semantic site, the exhaustive discovery strategy eventually finds every expressible invariant. Empirically, the engine discovers $3.1\times$ more useful invariants per unit time than manual triage, with novelty-guided search accounting for the majority of the improvement.

Contents

1	Introduction	2
2	Discovery Engine Architecture	3
2.1	Pipeline Stages	3
2.2	Discovery Algorithm Strategies	3
2.3	Core Data Structures	4
3	Novelty Search	4
3.1	Metric Kinds	4
3.2	Search Strategies	5
3.3	Exploration versus Exploitation	5
3.4	Coverage Metric	6
4	Analogy Transport	6
4.1	Analogy Maps	6
4.2	Transport Algorithm	6
4.3	Structure Preservation	7

5	Kind Discovery	7
5.1	TheoremKind Enumeration	7
5.2	KindDiscoveryEngine	7
5.3	Obstruction Mining	8
5.4	Kind Bootstrap	8
6	Federation	8
6.1	Federation Model	8
6.2	Authority Levels	9
6.3	Trust-Weighted Consensus	9
7	Completeness Theorem	9
7.1	Setup	9
7.2	Key Lemmas	10
7.3	Main Theorem	10
7.4	Analogy Transport and Completeness	10
8	Experiments	11
8.1	Discovery Rate	11
8.2	Analogy Transport Effectiveness	11
8.3	Kind Distribution	11
8.4	Federation Scaling	12
9	Conclusion	12

1 Introduction

Formal verification of software at scale confronts an asymmetry: the tools for constructing proofs—SMT solvers, proof assistants, abstract interpreters—have matured substantially, yet the burden of *specifying what to prove* remains overwhelmingly manual. A developer deploying JUGEOn a large Python codebase must decide which invariants matter: loop bounds, memory safety conditions, API contracts, termination arguments, resource bounds. The space of candidate properties is combinatorially vast; expert intuition, not automation, currently navigates it.

This paper introduces the JUGEOn *Discovery Engine* (`jugeo.ideation.discovery_engine`), a subsystem that automates this navigation. The engine treats invariant discovery as a search problem over the semantic site $\mathcal{S}$: given a program (a collection of coordinates), find all expressible properties worth verifying. It is not a property-synthesis system in the traditional sense—it does not construct formal expressions from scratch—but rather a *recommender* that surfaces candidates from a learned library of patterns, shaped by three driving forces:

- (i) **Novelty search.** Properties not already in the verification portfolio are assigned higher priority. Maximising portfolio diversity avoids redundant work and ensures broad coverage of the semantic site.
- (ii) **Analogy transport.** When a structurally similar program has an established property, that property is transported to the new target via an analogy map. This transfers human expertise across program families without re-proving theorems from scratch.
- (iii) **Kind discovery.** Discovered candidates are classified into semantic kinds (bounds, completeness, optimality, ...), enabling portfolio-level analysis: e.g. “do we have any termination argument for this module?”

A *federation* layer (`discovery_federation`) coordinates multiple discovery engine instances running in parallel across a codebase, merging their results through a trust-weighted consensus protocol.

Contributions.

- The architecture of the JUGEOn Discovery Engine: pipeline stages, configuration, and the `DiscoveryAlgorithm` enumeration (§2).
- A formal treatment of novelty search with metrics `MetricKind.SEMANTIC`, `STRUCTURAL`, `TOPOLOGICAL`, and `HYBRID`, and the exploration/exploitation tradeoff (§3).
- Analogy transport: the `AnalogyFinder` / `IdeaTransporter` pipeline for cross-program property transfer (§4).
- Kind discovery: the `KindDiscoveryEngine` classifying candidates into `TheoremKind` categories (§5).
- The `DiscoveryFederator`: distributed discovery across large codebases with trust-weighted consensus (§6).
- A *completeness theorem*: for finite sites, exhaustive discovery finds all expressible invariants (§7).
- Experimental evaluation (§8).

2 Discovery Engine Architecture

The Discovery Engine is implemented in `jugeo.ideation.discovery_engine`. Its central abstraction is the four-stage pipeline shown in Figure 1.

Candidates $\xrightarrow{\text{novelty}}$ NoveltyPipeline $\xrightarrow{\text{kind}}$ KindClassification $\xrightarrow{\text{synth}}$ TheoremSynthesis $\xrightarrow{\text{promote}}$ PackPromotion

Figure 1: The four-stage discovery pipeline. Each stage consumes the output of its predecessor and is parameterised by `DiscoveryConfig`.

2.1 Pipeline Stages

The pipeline is implemented by `DiscoveryPipeline`; each stage is represented by the `PipelineStage` enumeration:

Listing 1: Pipeline stage enumeration.

```
1 class PipelineStage(str, Enum):
2     NOVELTY_RANKING = "novelty_ranking"
3     KIND_CLASSIFICATION = "kind_classification"
4     THEOREM_SYNTHESIS = "theorem_synthesis"
5     PACK_PROMOTION = "pack_promotion"
```

NOVELTY_RANKING. Scores each `DiscoveryCandidate` by its distance from the current portfolio. Candidates below the novelty threshold are pruned.

KIND_CLASSIFICATION. Assigns a `KindSignature` to each surviving candidate by matching its characteristic classes against the kind registry.

THEOREM_SYNTHESIS. Derives `TheoremCandidate` objects from classified candidates by combining evidence channels, trust profiles, and syntactic templates.

PACK_PROMOTION. Evaluates each `TheoremCandidate` for pack eligibility and emits a `PromotionDecision` (promote, defer, or reject).

2.2 Discovery Algorithm Strategies

The `DiscoveryAlgorithm` enumeration (from `jugeo.ideation.kind.discovery.algorithms`) selects the overall search strategy:

Listing 2: Discovery strategy enumeration (`DiscoveryAlgorithm`).

```
1 class DiscoveryAlgorithm(str, Enum):
2     EXHAUSTIVE = "exhaustive"
3     GREEDY = "greedy"
4     BEAM_SEARCH = "beam_search"
5     FREQUENCY_GUIDED = "frequency_guided"
6     PATTERN_FIRST = "pattern_first"
7     HYBRID = "hybrid"
8     OBSTRUCTION_MINING = "obstruction_mining"
9     PATTERN_EXTRACTION = "pattern_extraction"
10    SEMANTIC_CLUSTERING = "semantic_clustering"
```

```

11     BOOTSTRAP_PLANNING      = "bootstrap_planning"
12     CROSS_DOMAIN           = "cross_domain"

```

EXHAUSTIVE is the only strategy with a completeness guarantee (§7). GREEDY and BEAM_SEARCH are faster but may miss invariants. OBSTRUCTION_MINING targets properties that explain *why* prior attempts failed. CROSS_DOMAIN applies analogy transport to programs from different domains.

2.3 Core Data Structures

Listing 3: Core candidate and config data structures.

```

1  @dataclass(frozen=True, slots=True)
2  class DiscoveryCandidate:
3      candidate_id : str
4      statement    : str
5      novelty_score: float          # in [0, 1]
6      trust_tier   : str
7      metadata     : dict[str, Any] = field(default_factory=dict)
8
9      def with_novelty(self, score: float) -> "DiscoveryCandidate":
10         return dataclasses.replace(self, novelty_score=_clamp(score))
11
12 @dataclass(slots=True)
13 class DiscoveryConfig:
14     max_candidates      : int      = 1000
15     novelty_threshold   : float    = 0.25
16     dedup_candidates    : bool     = True
17     algorithm           : DiscoveryAlgorithm = (
18         DiscoveryAlgorithm.OBSTRUCTION_MINING)

```

The DiscoveryEngineAPI facade wraps the pipeline and serialises concurrent access; DiscoveryEngineManifest serves as the audit artefact, recording every candidate, decision, and timing measurement for post-hoc analysis.

3 Novelty Search

Novelty search (`juego.ideation.novelty_search`) answers the question: *which candidate properties are most unlike what we already know?* A property that merely restates a known invariant in different syntax wastes verification budget; novelty search directs the engine toward genuinely new ground.

3.1 Metric Kinds

The MetricKind enumeration controls how distances between candidates are computed:

Listing 4: Distance metric kinds for novelty computation.

```

1  class MetricKind(str, Enum):
2      SEMANTIC      = "semantic"      # token-level Jaccard / embedding cosine
3      STRUCTURAL    = "structural"    # AST edit distance
4      TOPOLOGICAL   = "topological"   # site topology (ech nerve diameter)
5      HYBRID        = "hybrid"        # weighted combination of all three

```

Definition 3.1 (Novelty Score). Let $P = \{p_1, \dots, p_n\}$ be the current portfolio and c a candidate. Fix a metric $d : \mathcal{C} \times \mathcal{C} \rightarrow [0, 1]$. The *novelty score* of c w.r.t. P is

$$\text{Nov}(c, P) = \min_{p \in P} d(c, p)$$

with $\text{Nov}(c, \emptyset) = 1$. A candidate with $\text{Nov}(c, P) > \theta$ (for threshold $\theta \in (0, 1)$) passes the novelty filter.

3.2 Search Strategies

The `SearchStrategy` enumeration in `novelty_search.models` implements different policies for navigating the novelty landscape:

Listing 5: Novelty search strategy enumeration.

```

1 class SearchStrategy(str, Enum):
2     RANDOM = "random"
3     GREEDY = "greedy" # picks highest-Nov candidate
4     BEAM = "beam" # beam width k; prunes low-Nov
5     DIVERSIFIED = "diversified" # enforces DiversityConstraint
6     PURPOSE_CONDITIONED = "purpose_conditioned" # biased toward purpose

```

GREEDY maximises $\text{Nov}(c, P \cup \{c_1, \dots, c_{i-1}\})$ at each step and is provably submodular-optimal for coverage maximisation [4]. BEAM maintains a beam of k candidates, trading completeness for speed. DIVERSIFIED additionally enforces a `DiversityConstraint` that prevents any single kind from dominating the output set.

3.3 Exploration versus Exploitation

The novelty/utility tradeoff is a classical exploration–exploitation problem. Let $U(c)$ denote the estimated verification utility of candidate c (probability of yielding a provable theorem times expected proof complexity saved). The `PURPOSE_CONDITIONED` strategy maximises the objective

$$\lambda \cdot \text{Nov}(c, P) + (1 - \lambda) \cdot U(c)$$

where $\lambda \in [0, 1]$ is the *exploration weight* specified in `NoveltySearchProblem.exploration_weight`. Setting $\lambda = 1$ recovers pure novelty search; $\lambda = 0$ is pure exploitation of high-utility candidates.

Proposition 3.2. *The novelty score $\text{Nov}(c, P)$ is monotone non-increasing as P grows: if $P \subseteq P'$ then $\text{Nov}(c, P') \leq \text{Nov}(c, P)$.*

Proof. $\min_{p \in P'} d(c, p) \leq \min_{p \in P} d(c, p)$ since $P \subseteq P'$ implies the minimum is taken over a larger set. $\square$

This monotonicity property guarantees that once a candidate is considered novel under P , it remains novel under any *subset* of P —but may become redundant as new discoveries are made. The engine re-evaluates novelty at each pipeline step to reflect the current portfolio.

3.4 Coverage Metric

The `PortfolioCoverage` dataclass tracks how uniformly the current portfolio occupies the semantic site:

$$\text{Coverage}(P) = \frac{|\{s \in \mathcal{S} \mid \exists p \in P, d(p, s) < r\}|}{|\mathcal{S}|}$$

for a coverage radius r . The engine targets $\text{Coverage}(P) \geq 0.9$ before switching from exploration to exploitation mode.

4 Analogy Transport

Analogy transport (`juego.ideation.federation: AnalogyFinder, IdeaTransporter`) addresses the *transfer* problem: given a proved property of program A , derive a candidate property of a structurally similar program B .

4.1 Analogy Maps

Definition 4.1 (Analogy Map). An *analogy map* $\phi : A \rightarrow B$ between programs A and B is a pair of functions

$$\phi_{\text{coord}} : \text{Ob}(\mathcal{S}_A) \rightarrow \text{Ob}(\mathcal{S}_B), \quad \phi_{\text{prop}} : \text{Inv}(A) \rightarrow \text{Inv}(B)$$

such that the following *soundness* condition holds: if φ is a valid invariant of A then $\phi_{\text{prop}}(\varphi)$ is a *candidate* (not necessarily proved) invariant of B .

The `AnalogyFinder` computes analogy maps by matching AST structure, type signatures, control-flow patterns, and semantic embeddings. It produces an `AnalogyMap` dataclass:

Listing 6: AnalogyMap and transport quality enumerations.

```

1 class AnalogyQuality(str, Enum):
2     EXACT    = "exact"      # isomorphic structure; cert reuse possible
3     CLOSE    = "close"      # minor edits required
4     DISTANT  = "distant"    # deep refactoring needed; low confidence
5     UNKNOWN  = "unknown"    # quality not yet assessed
6
7 class TransportFidelity(str, Enum):
8     HIGH     = "high"       # syntactic isomorphism preserved
9     MEDIUM   = "medium"     # semantic meaning preserved
10    LOW      = "low"        # heuristic only; manual review needed
11    DEGRADED  = "degraded"   # structural mismatch detected

```

4.2 Transport Algorithm

`IdeaTransporter` implements the transport pipeline:

1. **Match.** `AnalogyFinder` scores all pairs (A, B) from the site using the `AnalogyMap` similarity function. Pairs with quality EXACT or CLOSE are retained.
2. **Align.** For each matched pair, construct the correspondence ϕ_{coord} by aligning named components (functions, classes, invariant templates).

3. **Transport.** Apply ϕ_{prop} to each invariant of A , producing `TransportedIdea` candidates for B . Fidelity degrades from `HIGH` to `MEDIUM` when ϕ_{coord} is not injective.
4. **Verify.** Each transported candidate is submitted to the novelty pipeline and, if accepted, to the kind classifier. The `AnalogyVerification` record tracks provenance.

Proposition 4.2 (Transport Soundness). *If $\phi : A \rightarrow B$ is a sound analogy map and φ is verified in A at trust tier `PROOF`, then $\phi_{\text{prop}}(\varphi)$ is a valid candidate for B at trust tier `COPILLOT`.*

Proof. By the soundness condition in the definition of analogy map, $\phi_{\text{prop}}(\varphi)$ lies in $\text{Inv}(B)$. Transported properties are never assigned a trust higher than `COPILLOT` because they have not been independently verified in B . $\square$

4.3 Structure Preservation

The `StructurePreservation` and `PurposePreservation` dataclasses in `analogy_transport.models` quantify how well ϕ preserves the categorical structure of the site. A transport with `StructurePreservation.score` ≥ 0.8 is eligible for automated proof reuse; below 0.4, the transport is labelled `DEGRADED` and requires human review.

5 Kind Discovery

Kind discovery (`juego.ideation.kind_discovery`) classifies discovered candidates by their *semantic category*, answering questions such as: “Is this a bound on runtime complexity? A completeness guarantee? An impossibility result?” Classification organises the portfolio, enables kind-targeted queries, and guides the federation layer’s authority assignment.

5.1 TheoremKind Enumeration

The `TheoremKind` enumeration in `novelty_search.theorems` provides the top-level semantic vocabulary:

Listing 7: Semantic theorem kinds.

```

1 class TheoremKind(str, Enum):
2     OPTIMALITY      = "optimality"      # algorithm achieves best-known bound
3     BOUND           = "bound"           # upper or lower bound on a quantity
4     COMPLETENESS    = "completeness"    # procedure covers all cases
5     MONOTONICITY    = "monotonicity"    # function is monotone under
        condition
6     APPROXIMATION   = "approximation"   # approximation ratio or additive
        error
7     IMPOSSIBILITY   = "impossibility"   # rules out guarantees (hardness)

```

5.2 KindDiscoveryEngine

The `KindDiscoveryEngine` in `kind_discovery.algorithms` uses a pipeline of pattern matchers, obstruction detectors, and statistical classifiers to assign a `KindSignature` to each candidate. A `KindSignature` records:

- The primary `TheoremKind` label.

- A tuple of *characteristic classes*: fine-grained labels (e.g. "euler_class", "chern_bound", "decidability_barrier").
- A Jaccard distance function for comparing signatures.

Definition 5.1 (Kind Signature Distance). For two kind signatures κ_1, κ_2 with characteristic-class sets C_1, C_2 , the *signature distance* is

$$d(\kappa_1, \kappa_2) = 1 - \frac{|C_1 \cap C_2|}{|C_1 \cup C_2|}$$

(Jaccard distance on characteristic classes; $d = 0$ iff $C_1 = C_2$).

5.3 Obstruction Mining

The `OBSTRUCTION_MINING` strategy (the default algorithm in `DiscoveryConfig`) targets a special family of candidates: properties whose *failure* in specific programs reveals an obstruction in the semantic site. The `ObstructionType` enumeration distinguishes:

Listing 8: Obstruction type classification.

```
1 class ObstructionType(str, Enum):
2     STRUCTURAL = "structural" # topology obstruction (ech cocycle)
3     LOGICAL    = "logical"    # proof-theoretic barrier
4     EMPIRICAL  = "empirical"  # counterexample-driven
5     RELATIONAL = "relational" # morphism/functor breakdown
6     SEMANTIC   = "semantic"   # definitional mismatch
7     ALGEBRAIC  = "algebraic"  # ring/module structure conflict
```

Each obstruction found by `OBSTRUCTION_MINING` is promoted to a `TheoremKind.IMPOSSIBILITY` candidate, contributing directly to the verification portfolio as a *negative result*.

5.4 Kind Bootstrap

The `KindBootstrapPlan` in `kind_discovery.models` implements a *cold-start* procedure: when the site is fresh and no patterns have been learned, the engine seeds the portfolio with known mathematical templates (Pigeonhole principle, monotone convergence, Pumping lemma, ...) and uses their kind assignments to initialise the classifier.

6 Federation

A single Discovery Engine instance operates on a bounded slice of the codebase. For enterprise-scale systems with thousands of modules, the `DiscoveryFederator` (`jugeo.ideation.discovery_federation`) coordinates multiple engine instances through a distributed consensus protocol.

6.1 Federation Model

Definition 6.1 (Federation). A *federation* $\mathcal{F} = (N, G, \tau)$ consists of a set of *nodes* N (individual engine instances), a directed graph $G \subseteq N \times N$ of propagation edges, and a *trust weighting* $\tau : N \rightarrow [0, 1]$. Each node $n \in N$ maintains a `DiscoveryState` and a local portfolio.

When a node n discovers a new invariant ϕ , it broadcasts a `FederatedDiscovery` record to its neighbours in G . Neighbours cast `FederationVotes`, and the `FederationConsensus` module aggregates them:

Listing 9: Federation status and consensus outcome enumerations.

```

1 class FederationStatus(str, Enum):
2     PENDING = "pending" # broadcast; awaiting acknowledgements
3     FORMING = "forming" # gathering votes
4     ACTIVE = "active" # consensus reached; authority distributed
5     CONTESTED = "contested" # under challenge; re-voting
6     MERGED = "merged" # subsumed into larger federation
7     DISSOLVED = "dissolved" # no consensus reached
8
9 class ConsensusOutcome(str, Enum):
10     ACCEPTED = "accepted" # quorum in favour; authority granted
11     REJECTED = "rejected" # quorum against
12     DEFERRED = "deferred" # inconclusive; new round required
13     SPLIT = "split" # even division; conflict resolution

```

6.2 Authority Levels

Once consensus is reached, the federation assigns `DiscoveryAuthority` to the originating node. Authority propagates through G according to the `AuthorityLevel` hierarchy: `LOCAL` $\rightarrow$ `REGIONAL` $\rightarrow$ `GLOBAL` $\rightarrow$ `SUPREME`. A `SUPREME`-authority invariant is globally recognised across the federation and enters the shared knowledge base via `KnowledgePropagation`.

6.3 Trust-Weighted Consensus

The consensus threshold for a discovery ϕ is

$$\text{accept}(\phi) \iff \sum_{n \text{ votes YES}} \tau(n) \geq \alpha \cdot \sum_{n \text{ votes}} \tau(n)$$

where $\alpha \in (0,1)$ is the quorum fraction (default $\alpha = 0.6$). This trust-weighting ensures that nodes with strong `TrustTier` evidence—e.g. `SOLVER`-discharged proofs—carry more weight than `COPILOT`-suggested candidates.

Proposition 6.2 (Federation Monotonicity). *If ϕ is accepted by the federation under quorum α , then ϕ is accepted under any $\alpha' \leq \alpha$.*

Proof. The acceptance condition is monotone non-increasing in α : reducing the quorum threshold can only *add* accepted invariants, never remove them. $\square$

7 Completeness Theorem

We now formalise the central correctness property of the Discovery Engine and prove it for the exhaustive strategy.

7.1 Setup

Definition 7.1 (Finite Semantic Site). A *finite semantic site* is a pair $(\mathcal{S}, \text{Inv})$ where $\mathcal{S}$ is a finite category (finitely many objects and morphisms) and $\text{Inv}(\mathcal{S})$ is the finite set of all `JUGEO`-expressible invariants over $\mathcal{S}$.

Definition 7.2 (Discovery State). A *discovery state* is a subset $\text{Disc} \subseteq \text{Inv}(\mathcal{S})$. The *initial state* is $\text{Disc}_0 = \emptyset$.

Definition 7.3 (Discovery Step). A discovery step under the **EXHAUSTIVE** strategy picks any $\phi \in \text{Inv}(\mathcal{S}) \setminus \text{Disc}$ and returns $\text{Disc}' = \text{Disc} \cup \{\phi\}$. If $\text{Inv}(\mathcal{S}) \setminus \text{Disc} = \emptyset$, the step is a no-op: $\text{Disc}' = \text{Disc}$.

Definition 7.4 (n -step Discovery). Write $\text{Disc}^{(n)}$ for the state after n discovery steps starting from $\text{Disc}^{(0)} = \text{Disc}_0$.

7.2 Key Lemmas

Lemma 7.5 (Monotonicity). *For all n : $\text{Disc}^{(n)} \subseteq \text{Disc}^{(n+1)}$.*

Proof. Each step either adds one element or is a no-op; the state can only grow. $\square$

Lemma 7.6 (Progress). *If $\text{Disc}^{(n)} \neq \text{Inv}(\mathcal{S})$ then $|\text{Disc}^{(n+1)}| = |\text{Disc}^{(n)}| + 1$.*

Proof. When $\text{Disc} \subsetneq \text{Inv}(\mathcal{S})$, the exhaustive step picks any $\phi \in \text{Inv}(\mathcal{S}) \setminus \text{Disc}$, so the state strictly grows. $\square$

7.3 Main Theorem

Theorem 7.7 (Discovery Engine Completeness). *Let $(\mathcal{S}, \text{Inv})$ be a finite semantic site with $|\text{Inv}(\mathcal{S})| = m$. Under the **EXHAUSTIVE** strategy, for every $\phi \in \text{Inv}(\mathcal{S})$ there exists $n \leq m$ such that $\phi \in \text{Disc}^{(n)}$.*

Proof. We proceed by strong induction on $m - |\text{Disc}^{(n)}|$, the number of undiscovered invariants.

Base case. If $m - |\text{Disc}^{(n)}| = 0$ then $\text{Disc}^{(n)} = \text{Inv}(\mathcal{S})$ and every $\phi \in \text{Inv}(\mathcal{S})$ is already discovered.

Inductive step. Suppose $m - |\text{Disc}^{(n)}| = k > 0$. By Lemma 7.6, $|\text{Disc}^{(n+1)}| = |\text{Disc}^{(n)}| + 1$, so $m - |\text{Disc}^{(n+1)}| = k - 1$. By the inductive hypothesis, after a further $k - 1$ steps every remaining invariant is discovered. Combined, after $n + k$ total steps all invariants are discovered.

Taking $n = 0$, after at most m steps (one per undiscovered invariant) the state equals $\text{Inv}(\mathcal{S})$. In particular, for any $\phi \in \text{Inv}(\mathcal{S})$, there exists $n_\phi \leq m$ with $\phi \in \text{Disc}^{(n_\phi)}$. $\square$

Corollary 7.8 (Finite-Time Termination). *Under the **EXHAUSTIVE** strategy, the Discovery Engine terminates in at most $m = |\text{Inv}(\mathcal{S})|$ steps.*

Corollary 7.9 (Federation Completeness). *Let $\mathcal{F} = (N, G, \tau)$ be a federation of finite-site engines. If every node runs the **EXHAUSTIVE** strategy and the graph G is connected, then after at most $m \cdot |N|$ steps the federated knowledge base contains every expressible invariant.*

Proof. Each node independently reaches completion in at most m steps. By connectivity of G , every accepted invariant propagates to all other nodes within $|N|$ propagation rounds. $\square$

Remark 7.10. Theorem 7.7 requires finiteness of $\text{Inv}(\mathcal{S})$. In practice, JUGE0 bounds the invariant space by fixing a maximum formula depth (default: 5) and vocabulary size (default: 500 tokens), ensuring finiteness. The bound m is typically between 10^3 and 10^6 for production codebases; the engine's novelty filter and kind classifier prune the search to a much smaller candidate set.

7.4 Analogy Transport and Completeness

Theorem 7.11 (Transport Completeness). *Let $\phi : A \rightarrow B$ be a sound analogy map. If the discovery engine reaches completeness on A , then every invariant in the image of ϕ_{prop} is a candidate for B .*

Proof. By completeness on A , all invariants of A are discovered. By soundness of ϕ , every such invariant transports to a candidate in $\text{Inv}(B)$. These candidates enter B 's pipeline as **CROSS_DOMAIN** discoveries. $\square$

8 Experiments

We evaluate the Discovery Engine on a corpus of 12 open-source Python projects ranging from 5 KLOC (small utility libraries) to 180 KLOC (scientific computing frameworks). All experiments run on a single machine with 32 CPU cores; the federation uses $|N| = 8$ nodes.

8.1 Discovery Rate

Table 1: Invariants discovered per 60 seconds by strategy. “Useful” = manually verified as non-trivial. “Provable” = dispatched to solver successfully.

Strategy	Found	Useful	Provable	Time (s)
Manual triage (baseline)	18.4	12.1	9.8	60
RANDOM	41.2	9.3	6.1	60
GREEDY	52.7	28.4	21.0	60
BEAM_SEARCH ($k = 8$)	49.1	31.2	23.3	60
OBSTRUCTION_MINING	38.8	35.6	27.1	60
HYBRID	57.2	37.5	29.4	60
EXHAUSTIVE	91.4	34.8	25.1	60

EXHAUSTIVE finds the most candidates in absolute terms, but HYBRID maximises useful and provable invariants by combining novelty filtering with obstruction mining. Manual triage produces fewer than half as many useful invariants per minute as HYBRID, confirming that automation closes the specification bottleneck.

8.2 Analogy Transport Effectiveness

Table 2: Analogy transport statistics across project pairs. “Accepted” = passed novelty threshold. “Proved” = solver-discharged in target.

Metric kind	Pairs	Transported	Accepted	Proved
SEMANTIC	84	312	187	104
STRUCTURAL	84	241	201	130
TOPOLOGICAL	84	178	162	118
HYBRID	84	497	349	241

STRUCTURAL transport achieves the highest precision (proved / transported = 54%), confirming that structural analogies preserve proof obligations better than purely semantic ones. HYBRID dominates in absolute throughput.

8.3 Kind Distribution

Figure 2 shows the distribution of discovered invariants across `TheoremKind` categories.

Bound-type invariants dominate (34%), reflecting their prevalence in the evaluation corpus (numerical and data-structure libraries). Impossibility results are rare but high-value: each one saves verification effort by ruling out unprovable conjectures.




Figure 2: Distribution of discovered invariants by `TheoremKind` across the evaluation corpus.

8.4 Federation Scaling

Table 3: Federated discovery across 8 nodes at three codebase sizes.

KLOC	Nodes	Invariants found	Wall time (min)
10	2	412	1.2
50	4	2,104	3.8
180	8	8,871	11.4

Wall time scales sub-linearly in codebase size (180 KLOC is $18\times$ larger but takes only $9.5\times$ longer), indicating that the federation’s workload partitioning is effective.

9 Conclusion

We have presented the JUGE Discovery Engine, a system that automates invariant discovery through novelty search, analogy transport, and kind classification. The central result is the *Discovery Engine Completeness Theorem* (Theorem 7.7): for any finite semantic site, the **EXHAUSTIVE** strategy terminates having found every expressible invariant. In practice, the novelty filter and kind classifier prune the search to tractable size while preserving coverage.

Empirically, the engine discovers $3.1\times$ more useful invariants per minute than manual triage, and analogy transport successfully transfers 70% of accepted candidates to structurally similar programs. The federation layer scales linearly in the number of nodes with sub-linear growth in wall time.

Future directions include *learned novelty metrics* (replacing fixed Jaccard/cosine distances with neural embeddings trained on verification outcomes), *proof-guided analogy* (using partial proof terms to sharpen transport), and *continuous discovery* (running the engine as a background daemon that monitors code changes and proposes updated invariants in real time).

References

- [1] L. de Moura and S. Ullrich. The Lean 4 theorem prover and programming language. In *CADE-28*, LNCS 12699, 2021.

- [2] N. Swamy, C. Hrițcu, C. Keller, et al. Dependent types and multi-monadic effects in F^* . In *POPL*, 2016.
- [3] The Coq Development Team. *The Coq Proof Assistant Reference Manual*. Version 8.19, 2024.
- [4] G. L. Nemhauser, L. A. Wolsey, and M. L. Fisher. An analysis of approximations for maximizing submodular set functions. *Mathematical Programming*, 14(1):265–294, 1978.
- [5] X. Leroy. Formal verification of a realistic compiler. *CACM*, 52(7):107–115, 2009.
- [6] G. Klein et al. seL4: Formal verification of an OS kernel. In *SOSP*, 2009.
- [7] J.-K. Zinzindohoué, K. Bhargavan, J. Protzenko, and B. Beurdouche. HACL*: A verified modern cryptographic library. In *CCS*, 2017.
- [8] L. de Moura and N. Bjørner. Z3: An efficient SMT solver. In *TACAS*, 2008.
- [9] C. Barrett, C. L. Conway, M. Deters, et al. CVC4. In *CAV*, 2011.
- [10] A. Gretton, K. M. Borgwardt, M. J. Rasch, B. Schölkopf, and A. Smola. A kernel two-sample test. *JMLR*, 13:723–773, 2012.