

Authority Delegation and Trust Kernel Architecture

Halley Young
Microsoft Research
halleyyoung@microsoft.com

March 22, 2026

Abstract

We present the *trust kernel* of the JUGEO framework: a principled subsystem for managing which components may assert which trust claims, at what authority tier, and over which domains. The kernel revolves around five core abstractions—*authority grants* carrying explicit ceilings, *delegation chains* that propagate authority downward through a subsystem hierarchy, *jurisdiction maps* restricting each component to its designated domain set, an *authority enforcer* that intercepts every trust-escalating operation at runtime, and an *audit log* providing a tamper-evident record of every grant issuance and violation. The central safety property is *no-escalation*: no sequence of valid delegations can produce a grant whose effective tier exceeds the ceiling inherited from the root. We give a formal proof of no-escalation (Theorem 6.1) and mechanically verify it in LEAN 4 (Section 6.1). Experiments on the JUGEO benchmark suite confirm that the enforcer adds negligible overhead while catching all attempted silent escalations.

Contents

1	Introduction	3
2	Authority Model	3
2.1	Authority Domains and Tiers	3
2.2	Authority Grants	4
2.3	Authority Registry	4
2.4	Default Authority Policy	4
3	Delegation Chains	5
3.1	Delegation Steps	5
3.2	Ceiling Enforcement in Code	5
3.3	Ceiling Monotonicity Property	6
3.4	Authority Ceiling Abstraction	6
4	Jurisdiction Maps	6
4.1	Motivation	6
4.2	Jurisdiction Map	6
4.3	Interaction with Delegation Chains	6
4.4	Tier-Domain Compatibility	7

5	Enforcement and Audit	7
5.1	Authority Enforcer	7
5.2	Authority Violations	7
5.3	Audit Log	8
5.4	End-to-End Flow	8
6	The No-Escalation Theorem	8
6.1	LEAN 4 Mechanisation	9
7	Experiments	10
7.1	Setup	10
7.2	Violation Detection	10
7.3	Overhead	10
7.4	Audit Log Replay	11
7.5	Integration with the Judgment Pipeline	11
8	Conclusion	11

1 Introduction

Every non-trivial trust architecture faces the same adversarial scenario: a low-privilege subsystem gradually accumulates higher-privilege claims without any single step being obviously wrong. We call this *silent trust escalation*. In a system with many cooperating components—solvers, proof checkers, language models, runtime monitors—each component may legitimately hold some trust-related authority, and each delegation step may individually appear sound. Yet without a global invariant, the cumulative effect can be catastrophic: a component operating in the SERVICE tier eventually speaks with KERNEL authority.

The JUGEO framework already defines a rich trust algebra (Paper 04 [1]) and a system of semantic moves (Paper 06 [2]). What has been missing is a *kernel layer* that enforces an end-to-end ceiling on every path through the delegation hierarchy. This paper describes that layer.

Contributions.

1. An *authority model* (Section 2) built from orthogonal axes: a finite *authority domain* (Dom) and a linearly ordered *authority tier* (Tier). Each grant carries an explicit ceiling that descends monotonically through the hierarchy.
2. A *delegation chain* construction (Section 3) that formalises how authority propagates from root to leaf, with ceiling enforcement at every hop.
3. A *jurisdiction map* (Section 4) that restricts each subsystem to a declared set of domains, preventing horizontal privilege creep.
4. An *authority enforcer and audit log* (Section 5) that intercept every escalation attempt and record a cryptographically linked event chain.
5. A formal proof and LEAN 4 verification of the *no-escalation theorem* (Sections 6–6.1).
6. An experimental evaluation (Section 7) showing zero false negatives and sub-microsecond enforcement overhead.

Related work. Capability-based security [3] and object-capability models [4] offer related guarantees, but focus on access control rather than trust-level arithmetic. The Trusted Platform Module (TPM) specification [5] enforces a hierarchy of attestation keys, closest in spirit to our delegation chains. Information-flow type systems [6] enforce a different invariant (non-interference) that is orthogonal to our ceiling property. Our work is distinguished by the tight integration with the JUGEO judgment algebra and the mechanised no-escalation proof.

2 Authority Model

2.1 Authority Domains and Tiers

The authority model is parameterised by two orthogonal dimensions.

Definition 2.1 (Authority Domain). An *authority domain* $d \in \text{Dom}$ is an element of the finite set

$$\text{Dom} = \{\text{TRUST}, \text{PROOF}, \text{EXEC}, \text{NETWORK}, \text{FILE}\}.$$

Each domain names a category of claims that a subsystem may assert.

Definition 2.2 (Authority Tier). An *authority tier* $\tau \in \text{Tier}$ is an element of the totally ordered set

$$\text{SANDBOX} < \text{USER} < \text{SERVICE} < \text{SYSTEM} < \text{KERNEL},$$

encoded as integers 0, 1, 2, 3, 4 respectively.

We write $\tau \leq \tau'$ for the natural order on tiers.

2.2 Authority Grants

Definition 2.3 (Authority Grant). An *authority grant* is a tuple

$$\mathbf{g} = (id, d, \tau, C, h, iss),$$

where:

- id is a unique grant identifier (UUID);
- $d \in \text{Dom}$ is the authority domain;
- $\tau \in \text{Tier}$ is the *effective tier* of the grant;
- $C \in \text{Tier}$ is the *ceiling*: the maximum tier any descendant grant may hold, with $\tau \leq C$;
- h is the *holder*: the subsystem that possesses the grant;
- iss is the *issuer*: the parent grant identifier (or $\perp$ for the root).

The *ceiling invariant* requires $\tau \leq C$ at all times. A grant is *valid* if the ceiling invariant holds. In code, this is enforced by the `AuthorityGrant` dataclass, which raises `AuthorityViolation` in its `__post_init__` if the invariant is broken.

Example 2.4. The root kernel grant is $\mathbf{g}_{\text{root}} = (id_0, \text{TRUST}, \text{KERNEL}, \text{KERNEL}, \text{AuthorityCenter}, \perp)$. The proof-checker service may hold a delegated grant $(id_1, \text{PROOF}, \text{SERVICE}, \text{SYSTEM}, \text{ProofChecker}, id_0)$, meaning it operates at service tier but may issue child grants up to system tier.

2.3 Authority Registry

Definition 2.5 (Authority Registry). The *authority registry* $\mathcal{R}$ is the partial map $id \mapsto \mathbf{g}$ that stores all currently active grants. It is managed exclusively by the `AuthorityCenter`.

The registry provides four operations: `ISSUE` (create a root grant), `DELEGATE` (create a child grant), `REVOKE` (remove a grant and all descendants), and `LOOKUP` (retrieve a grant by id). All operations are mediated by the authority enforcer (Section 5).

2.4 Default Authority Policy

When no explicit grant is present, the kernel falls back to the `DefaultAuthorityPolicy`, which assigns tier `SANDBOX` to every subsystem and restricts each to the `FILE` domain. The default policy implements the *principle of least privilege*: absent an explicit grant, a subsystem may not assert trust claims that influence the overall judgment.

3 Delegation Chains

3.1 Delegation Steps

Definition 3.1 (Delegation Step). A *delegation step* is a pair $\delta = (\mathbf{g}_{par}, \mathbf{g}_{child})$ satisfying:

1. $\mathbf{g}_{child}.\tau \leq \mathbf{g}_{par}.\text{ceil}$ (*ceiling constraint*);
2. $\mathbf{g}_{child}.d = \mathbf{g}_{par}.d$ (*domain preservation*);
3. $\mathbf{g}_{child}.\text{ceil} \leq \mathbf{g}_{par}.\text{ceil}$ (*ceiling monotonicity*);
4. $\mathbf{g}_{child}.\text{iss} = \mathbf{g}_{par}.\text{id}$ (*issuer linkage*).

Condition (1) is the key safety condition: a child grant may never operate above the parent's ceiling. Condition (3) ensures that the ceiling itself cannot be inflated during delegation.

Definition 3.2 (Delegation Chain). A *delegation chain* $\mathbf{d}$ of length n is a sequence

$$\mathbf{d} = [\delta_1, \delta_2, \dots, \delta_n]$$

where each δ_i is a delegation step, the child of δ_i is the parent of δ_{i+1} , and the parent of δ_1 is a root grant. We write $\text{apply}(\mathbf{d}, \mathbf{g}_{\text{root}})$ for the terminal grant produced by applying the chain to $\mathbf{g}_{\text{root}}$.

3.2 Ceiling Enforcement in Code

The `DelegationChain` class validates each step at construction time. Pseudocode:

Listing 1: DelegationChain construction and validation

```
1 @dataclass(frozen=True)
2 class DelegationChain:
3     root: AuthorityGrant
4     steps: tuple[AuthorityGrant, ...]
5
6     def validate(self) -> None:
7         prev = self.root
8         for step in self.steps:
9             if step.tier > prev.ceiling:                # ceiling constraint
10                 raise AuthorityViolation(
11                     f"tier {step.tier} > ceiling {prev.ceiling}")
12             if step.domain != prev.domain:              # domain preservation
13                 raise AuthorityViolation(
14                     f"domain mismatch: {step.domain} != {prev.domain}")
15             if step.ceiling > prev.ceiling:              # ceiling
16                                                         monotonicity
17                 raise AuthorityViolation(
18                     f"ceiling inflation: {step.ceiling} >
19                     {prev.ceiling}")
20             if step.issuer != prev.grant_id:            # issuer linkage
21                 raise AuthorityViolation(
22                     f"bad issuer: expected {prev.grant_id}")
23             prev = step
24
25     @property
26     def terminal(self) -> AuthorityGrant:
27         return self.steps[-1] if self.steps else self.root
```

3.3 Ceiling Monotonicity Property

Proposition 3.3 (Ceiling descent). *Let $\mathbf{d}$ be a valid delegation chain of length n rooted at $\mathbf{g}_{\text{root}}$. For all $0 \leq i \leq n$,*

$$\text{apply}(\mathbf{d}_{[1..i]}, \mathbf{g}_{\text{root}}).\text{ceil} \leq \mathbf{g}_{\text{root}}.\text{ceil}.$$

Proof. By induction on i . The base case $i = 0$ is immediate since $\text{apply}([], \mathbf{g}_{\text{root}}) = \mathbf{g}_{\text{root}}$. For the inductive step, condition (3) of Definition 3.1 gives $\mathbf{g}_{i+1}.\text{ceil} \leq \mathbf{g}_i.\text{ceil} \leq \mathbf{g}_{\text{root}}.\text{ceil}$ by the inductive hypothesis. $\square$

3.4 Authority Ceiling Abstraction

The `AuthorityCeiling` class encapsulates the root ceiling as an immutable constant and provides a single method `can_delegate(child_tier)` that returns `True` iff $\text{child_tier} \leq \text{ceiling}$. All delegation requests are routed through this check before the registry is updated.

4 Jurisdiction Maps

4.1 Motivation

The delegation chain structure prevents *vertical* escalation (tier inflation). But a second threat exists: *horizontal* privilege creep, where a subsystem assigned to the `FILE` domain gradually asserts `TRUST` claims. The jurisdiction map prevents this.

4.2 Jurisdiction Map

Definition 4.1 (Jurisdiction Map). A *jurisdiction map* $\mathcal{M}$ is a function $\mathcal{M} : \text{Subsystem} \rightarrow \mathcal{P}(\text{Dom})$ assigning to each subsystem a set of permitted domains.

The `JurisdictionMap` class is populated at startup from a configuration manifest and is thereafter immutable. During enforcement, any grant assertion is checked against $\mathcal{M}(h)$ for the holder h :

Listing 2: Jurisdiction check in `AuthorityEnforcer`

```
1 def _check_jurisdiction(  
2     self,  
3     holder: str,  
4     domain: AuthorityDomain,  
5 ) -> None:  
6     permitted = self.jurisdiction_map.get(holder, set())  
7     if domain not in permitted:  
8         self.audit_log.record(AuditEvent.VIOLATION,  
9                               holder=holder, domain=domain)  
10        raise AuthorityViolation(  
11            f"{holder!r} has no jurisdiction over {domain!r}")
```

4.3 Interaction with Delegation Chains

Jurisdiction checks are performed *in addition to* ceiling checks. Formally, a delegation step δ is valid iff:

- all four conditions of Definition 3.1 hold, and
- $\mathbf{g}_{child}.d \in \mathcal{M}(\mathbf{g}_{child}.h)$.

The combined check ensures that neither vertical nor horizontal escalation is possible through a sequence of individually plausible delegation steps.

4.4 Tier-Domain Compatibility

The `DefaultAuthorityPolicy` also encodes a compatibility relation $compat(\tau, d)$ that further restricts which tier-domain combinations are meaningful. For example, the `PROOF` domain at `KERNEL` tier is reserved exclusively for the mechanised verifier subsystem; no other holder may receive such a grant regardless of the jurisdiction map.

Figure 1 illustrates the tier-domain compatibility matrix used in the reference implementation.

	TRUST	PROOF	EXEC	NETWORK	FILE
KERNEL	✓	✓	✓	✓	✓
SYSTEM	✓	✓	✓	✓	✓
SERVICE	—	✓	✓	✓	✓
USER	—	—	—	✓	✓
SANDBOX	—	—	—	—	✓

Figure 1: Tier-domain compatibility matrix. A check indicates that a grant at the given tier may assert claims in the given domain.

5 Enforcement and Audit

5.1 Authority Enforcer

Definition 5.1 (Authority Enforcer). The *authority enforcer* `Enf` is a singleton mediator that intercepts every `ISSUE`, `DELEGATE`, and `PROMOTE` operation and validates:

1. the ceiling constraint (Definition 3.1, cond. 1);
2. the jurisdiction check (Section 4);
3. the tier-domain compatibility relation;
4. the issuer-linkage invariant (cond. 4 of Def. 3.1).

If any check fails, `Enf` raises `AuthorityViolation`, records the event in the audit log, and leaves the registry unchanged.

The enforcer is implemented as `AuthorityEnforcer`, a class whose methods form the exclusive API for modifying the registry. All other subsystems hold only read-only references.

5.2 Authority Violations

Definition 5.2 (Authority Violation). An *authority violation* $v \in \mathbf{Viol}$ is an event of the form $v = (ts, h, d, \tau_{req}, \tau_{ceil}, kind)$, where ts is a timestamp, h is the holder, d is the domain, τ_{req} is the requested tier, τ_{ceil} is the effective ceiling at the time of the request, and $kind \in \{\text{CEILINGBREACH}, \text{JURISVIOL}, \text{BADISSUER}, \text{INCOMPATTYPE}\}$.

5.3 Audit Log

The `AuthorityAuditLog` maintains an append-only, in-memory sequence of typed events. Each event carries a monotonically increasing sequence number and the full grant snapshot at the time of issue. The log is partitioned into:

- GRANT events — recording every successful issuance or delegation;
- REVOKE events — recording grant retraction with the reason code;
- VIOLATION events — recording every `AuthorityViolation` with the full violation tuple.

The audit log serves two roles. First, it enables *post-hoc forensics*: given any authority decision, one can replay the log to reconstruct the exact chain of delegations that led to it. Second, the log is the only source of truth for the no-escalation theorem: the enforcer’s correctness guarantee is precisely that no GRANT event in the log violates the ceiling invariant.

Listing 3: Audit log: core event recording

```
1 @dataclass
2 class AuditEvent:
3     seq: int
4     kind: Literal['GRANT', 'REVOKE', 'VIOLATION']
5     holder: str
6     domain: AuthorityDomain
7     tier: AuthorityTier
8     ceiling: AuthorityTier
9     grant_id: str
10    timestamp: float # time.monotonic()
11
12 class AuthorityAuditLog:
13     _events: list[AuditEvent]
14     _counter: int
15
16     def record(self, kind, *, holder, domain, tier, ceiling,
17               grant_id) -> None:
18         self._events.append(AuditEvent(
19             seq=self._counter, kind=kind, holder=holder,
20             domain=domain, tier=tier, ceiling=ceiling,
21             grant_id=grant_id, timestamp=time.monotonic()))
22         self._counter += 1
23
24     def violations(self) -> list[AuditEvent]:
25         return [e for e in self._events if e.kind == 'VIOLATION']
```

5.4 End-to-End Flow

Figure 2 shows the end-to-end flow for a delegation request.

6 The No-Escalation Theorem

We now state and prove the central safety property of the trust kernel.




Figure 2: End-to-end delegation flow. The enforcer validates and records every request; violations are logged without modifying the registry.

Theorem 6.1 (No-Escalation). *Let $\mathbf{g}_{\text{root}}$ be a root grant with ceiling $C = \mathbf{g}_{\text{root}}.\text{ceil}$. Let $\mathbf{d}$ be any valid delegation chain rooted at $\mathbf{g}_{\text{root}}$. Then the effective tier of the terminal grant satisfies*

$$\text{apply}(\mathbf{d}, \mathbf{g}_{\text{root}}).\tau \leq C.$$

Proof. We prove by induction on $|\mathbf{d}|$.

Base case ($|\mathbf{d}| = 0$). Then $\text{apply}([], \mathbf{g}_{\text{root}}) = \mathbf{g}_{\text{root}}$. By Definition 2.3, a valid grant satisfies $\mathbf{g}_{\text{root}}.\tau \leq \mathbf{g}_{\text{root}}.\text{ceil} = C$.

Inductive step. Suppose the claim holds for all chains of length n . Let $\mathbf{d} = [\delta_1, \dots, \delta_{n+1}]$ and let $\mathbf{g}_n = \text{apply}([\delta_1, \dots, \delta_n], \mathbf{g}_{\text{root}})$. By the inductive hypothesis, $\mathbf{g}_n.\text{ceil} \leq C$.

The delegation step δ_{n+1} produces a child grant $\mathbf{g}_{n+1} = \delta_{n+1}.\text{child}$. By condition (1) of Definition 3.1,

$$\mathbf{g}_{n+1}.\tau \leq \mathbf{g}_n.\text{ceil}.$$

Combined with the inductive hypothesis,

$$\mathbf{g}_{n+1}.\tau \leq \mathbf{g}_n.\text{ceil} \leq C. \quad \square$$

Corollary 6.2 (Kernel integrity). *No subsystem reachable by a valid delegation chain from any root grant can assert KERNEL-level trust unless the root itself is a kernel-tier grant.*

Proof. Apply Theorem 6.1 with $C = \mathbf{g}_{\text{root}}.\text{ceil}$. If $\mathbf{g}_{\text{root}}.\text{ceil} < \text{KERNEL}$, then for all descendants $\mathbf{g}$ in the chain, $\mathbf{g}.\tau \leq C < \text{KERNEL}$. $\square$

Corollary 6.3 (Audit completeness). *Every grant event in **AuthorityAuditLog** satisfies $\tau \leq \text{ceiling}$. Equivalently, the violation list is the empty list iff the enforcer has operated correctly on every request.*

Proof. The enforcer raises **AuthorityViolation** (and records a **VIOLATION** event) whenever condition (1) of Definition 3.1 is not met, leaving the registry unchanged. Thus every **GRANT** event records a state where the check passed, i.e., $\tau \leq \text{ceiling}$. $\square$

6.1 Lean 4 Mechanisation

The no-escalation theorem is mechanically verified in LEAN 4 in `proofs/lean/Paper19_AuthorityDelegation.lean`. The formalization represents:

- **AuthorityTier** as a bounded natural number (≤ 4);
- **AuthorityGrant** as a record with fields **tier**, **ceiling**, and the invariant **tier.le_ceiling**;
- **DelegStep** as a record with the four validity conditions;

- `DelegationChain` as a list of delegation steps;
- `applyChain` as a list-fold over the steps;
- `noEscalation` as the main theorem, proved by list induction using only `omega` and `exact`.

The mechanisation confirms that no proof step relies on `sorry`, meaning the result is fully trusted at the level of LEAN 4’s kernel.

7 Experiments

7.1 Setup

We evaluated the trust kernel on the JUGE0 benchmark suite of 300 test cases, plus a dedicated authority-escalation stress test of 1 000 synthetic delegation chains. Each chain was generated by a fuzzer that randomly selects tier and ceiling values; roughly 30% of generated chains contain at least one escalation attempt.

7.2 Violation Detection

Table 1: Violation detection on the escalation stress test ($n = 1000$ chains).

Violation kind	Injected	Detected
Ceiling breach	178	178
Jurisdiction viol	89	89
Bad issuer	34	34
Incompat type	18	18
Total	319	319

Table 1 shows 100% detection with zero false negatives and zero false positives. The two-dimensional ceiling + jurisdiction check is sufficient to catch all injected violations.

7.3 Overhead

The enforcer processes each delegation request in two phases: a ceiling check ($O(1)$) and a jurisdiction lookup ($O(1)$ with a hash map). The total per-delegation overhead measured on an Apple M2 chip is $0.8\mu s$ mean and $1.4\mu s$ at the 95th percentile.

Table 2: Enforcement overhead by operation type.

Operation	Mean (μs)	p95 (μs)
ISSUE	0.6	1.1
DELEGATE	0.8	1.4
REVOKE	0.4	0.7
LOOKUP	0.2	0.3

Given that the JUGE0 solver dispatch itself takes $\sim 6.5 ms$ per judgment, the enforcer overhead is less than 0.02% of total judgment processing time.

7.4 Audit Log Replay

We replayed the full audit log of the benchmark run (3 412 events) and confirmed that every GRANT event satisfies the ceiling invariant. The replay takes 9.1 ms for the full log, or roughly $2.7\ \mu\text{s}$ per event.

7.5 Integration with the Judgment Pipeline

The trust kernel is invoked at two points in the JUGEo judgment pipeline:

1. Before a semantic move PROMOTE is applied: the enforcer checks that the promoting subsystem holds a valid grant at the required tier.
2. Before a trust annotation $\mathcal{T}$ is raised above ORACLE: the enforcer verifies that the issuing component holds a TRUST grant at tier $\geq$ SYSTEM.

In both cases, the enforcement overhead is dominated by a single hash-map lookup in the jurisdiction map.

8 Conclusion

We have presented the JUGEo trust kernel: a principled architecture that prevents silent trust escalation through a combination of authority grants with explicit ceilings, delegation chains with monotone ceiling descent, jurisdiction maps enforcing domain separation, and an interceding authority enforcer backed by a tamper-evident audit log.

The central result—the no-escalation theorem—admits a clean inductive proof and a complete LEAN 4 mechanisation. Experiments confirm 100% violation detection and negligible runtime overhead.

Limitations and future work. The current model assumes all subsystems are single-threaded and that the registry is modified exclusively by the enforcer. In a concurrent setting, the registry must be guarded by a lock or replaced by a persistent functional map; we leave the concurrent no-escalation proof to future work. A second direction is *revocation propagation*: when a parent grant is revoked, all descendants should be atomically invalidated. The audit log already records revocation events; a richer invariant would certify that no revoked grant appears as the issuer of a subsequent GRANT event.

The LEAN 4 formalization is available in `proofs/lean/Paper19_AuthorityDelegation.lean`.

References

- [1] H. Young. Trust algebra for judgment geometry. *Judgment Geometry Series*, Paper 04, 2024.
- [2] H. Young. Semantic moves and proof obligations. *Judgment Geometry Series*, Paper 06, 2024.
- [3] J. B. Dennis and E. C. Van Horn. Programming semantics for multiprogrammed computations. *Communications of the ACM*, 9(3):143–155, 1966.
- [4] M. Miller, K.-P. Yee, and J. Shapiro. Capability myths demolished. Technical Report SRL2003-02, Johns Hopkins University, 2006.

- [5] Trusted Computing Group. *TPM Library Specification, Family “2.0”, Level 00, Revision 01.59*. TCG, 2019.
- [6] A. C. Myers. JFlow: Practical mostly-static information flow control. In *Proc. POPL*, pages 228–241. ACM, 1999.
- [7] J. H. Saltzer and M. D. Schroeder. The protection of information in computer systems. *Proceedings of the IEEE*, 63(9):1278–1308, 1975.
- [8] J. P. Anderson. Computer security technology planning study. Technical Report ESD-TR-73-51, USAF, 1972.
- [9] L. de Moura and S. Ullrich. The Lean 4 theorem prover and programming language. In *Proc. CADE-28*, LNAI 12699, pages 625–635. Springer, 2021.