

Semantic Caching and Incremental Verification in the Judgment Geometry Framework

Halley Young
Microsoft Research
halleyyoung@microsoft.com

March 22, 2026

Abstract

Verification is expensive. A single call to an SMT solver or a mechanised proof assistant may cost tens of milliseconds per judgement, and realistic software projects contain thousands of verification targets. Yet most code edits are local: a developer changes one function, and the vast majority of previously verified judgements remain correct. We formalise this intuition through *semantic caching*: a principled mechanism for storing and reusing verification results keyed by a pair $(h_{\text{code}}, h_{\text{spec}})$ that jointly identifies the program fragment and its specification.

We introduce the *semantic cache* $\mathcal{C}$ as a triple $(\mathcal{I}, \Pi, \text{Inv})$, where $\mathcal{I}$ maps keys to **CacheEntry** records, Π governs capacity management via one of four **EvictionStrategy** variants (LRU, LFU, FIFO, TTL), and Inv propagates staleness through a dependency graph using one of three **CascadeStrategy** modes. A **Checkpoint** records a consistent snapshot of the entire cache for serialisation and rollback. A **CacheWarmer** predicts which entries will be needed next and pre-computes them during idle cycles.

Our central result—the *Cache Correctness Theorem*—states that any cache entry that has not been marked stale by Inv faithfully reflects the current verification status of the corresponding code fragment. We formalise the theorem in **LEAN4** with a complete, sorry-free proof. Empirically, semantic caching reduces average verification latency by $5.9\times$ on a suite of 300 programs under a realistic edit-compile-verify loop, with 89% of verification calls served from cache.

Contents

1	Introduction	2
2	Cache Model	3
2.1	Keys and hashes	3
2.2	Cache entries	3
2.3	Cache index	4
2.4	Semantic cache	4
3	Cache Policy	4
3.1	Eviction strategies	4
3.2	The CachePolicy record	5

4	Cache Invalidation	5
4.1	Invalidation reasons	5
4.2	The CacheInvalidator	5
4.3	Dependency graph integration	6
5	Cascade Strategies	6
5.1	The CascadeStrategy enum	6
5.2	Choosing a strategy	6
6	Checkpointing	7
6.1	Motivation	7
6.2	The Checkpoint record	7
6.3	CheckpointBuilder	7
6.4	CacheSerializer	8
7	Cache Warming	8
7.1	Motivation	8
7.2	Prediction model	8
7.3	Trigger kinds	8
8	Cache Correctness Theorem	9
8.1	Setup	9
8.2	Main theorem	9
9	Experiments	10
9.1	Setup	10
9.2	Hit rate	10
9.3	Latency reduction	10
9.4	Warming effectiveness	11
9.5	Checkpoint overhead	11
9.6	Comparison with baseline	11
10	Conclusion	12

1 Introduction

“The art of doing mathematics consists in finding that special case which contains all the germs of generality.”

— David Hilbert

The Judgment Geometry (JUGEO) framework [1] represents program verification as navigation through a geometric space of *judgements*—8-tuples $(c, \varphi, \mathcal{A}, \mathcal{E}, \mathcal{O}, \mathcal{B}, \mathcal{T}, \Pi)$ that encode what is being verified, the evidence supporting it, and the trust level earned. As a program evolves under repeated edits, the geometry of its judgement space changes: some coordinates shift, some propositions are strengthened, and some trust annotations must be downgraded. In principle, every edit could require a full re-verification of all judgements reachable from the modified coordinate. In practice, most judgements are *unaffected* by a given edit.

The incremental verification problem. Let P be a program and P' the result of editing one function f . Standard verification re-checks every judgement in P' . Incremental verification identifies the *minimal affected set* $\text{aff}(f, P) \subseteq \text{Coords}(P)$ and re-verifies only those judgements. The challenge is twofold. First, the affected set must be computed *precisely*: over-approximation wastes effort; under-approximation yields unsound results. Second, previously verified judgements must be *stored* in a retrievable form so they can be replayed without calling the solver again.

Semantic caching as the solution. We answer both challenges with a *semantic cache*: a data structure that maps pairs $(h_{\text{code}}(f), h_{\text{spec}}(\varphi))$ to `CacheEntry` records encoding the verification result, the trust level attained, and the timestamp of computation. The key insight is that two functions with the same code hash and the same spec hash must have the same verification result, regardless of context. This is the *semantic* aspect: the cache key captures meaning, not syntax.

Scope and contributions. This paper makes the following contributions.

- (i) **Cache model** (§2): formal definition of `SemanticCache`, `CacheEntry`, and `CacheIndex` with their algebraic properties.
- (ii) **Cache policy** (§3): the `CachePolicy` type parametrised by `EvictionStrategy` and its capacity guarantees.
- (iii) **Invalidation** (§4): the `CacheInvalidator` and `InvalidationReason` taxonomy.
- (iv) **Cascade strategies** (§5): three `CascadeStrategy` modes and their propagation semantics.
- (v) **Checkpointing** (§6): `Checkpoint` and `CheckpointBuilder` for save/restore of consistent cache snapshots.
- (vi) **Cache warming** (§7): the `CacheWarmer` and predictive pre-computation.
- (vii) **Cache Correctness Theorem** (§8): a formal proof that non-invalidated entries are sound.
- (viii) **Experiments** (§9): $5.9\times$ median speedup, 89% hit rate.

Relationship to earlier papers. Paper 01 [1] introduced semantic sites; the cache index is a finite sub-presheaf of the verification presheaf over those sites. Paper 04 [2] established the trust algebra; each cache entry carries a `TrustLevel` from that algebra. Paper 06 [3] defined semantic moves; `VERIFY` is the primary consumer of cache entries, and `REPAIR` is the primary producer of invalidations.

2 Cache Model

2.1 Keys and hashes

A *cache key* is a pair

$$k = (\mathbf{h}_{\text{code}}(c), \mathbf{h}_{\text{spec}}(\varphi)) \in \mathbb{B}^{256} \times \mathbb{B}^{256}$$

where $\mathbf{h}_{\text{code}}(c)$ is the SHA-256 digest of the normalised source text of coordinate c , and $\mathbf{h}_{\text{spec}}(\varphi)$ is the SHA-256 digest of the canonical form of specification φ . Both hashes are computed on *normalised* representations: variable names are α -renamed, comments stripped, and whitespace canonicalised. This ensures that two semantically equivalent code fragments with different formatting share a key.

Definition 2.1 (Cache key). The *key space* is $\mathcal{K} = \mathbb{B}^{256} \times \mathbb{B}^{256}$ with equality $k = k'$ iff both components are equal. For a coordinate c and specification φ , the *canonical key* is $\mathbf{key}(c, \varphi) = (\mathbf{h}_{\text{code}}(c), \mathbf{h}_{\text{spec}}(\varphi))$.

Under the random oracle model, key collisions occur with probability 2^{-256} , which we treat as negligible.

2.2 Cache entries

Definition 2.2 (Cache entry). A *cache entry* is a record

$$e = (\text{key}, \text{result}, \text{trust}, \text{timestamp}, \text{hits}, \text{provenance})$$

where:

- $\text{key} \in \mathcal{K}$ identifies the cached judgement;
- $\text{result} \in \{\text{Valid}, \text{Invalid}, \text{Unknown}\}$ is the verification outcome;
- $\text{trust} \in \mathfrak{T}$ is the trust level at which the result was established;
- $\text{timestamp} \in \mathbb{N}$ is a logical clock value recording when the entry was written;
- $\text{hits} \in \mathbb{N}$ counts how many times the entry has been served to callers;
- provenance is an optional reference to the `Checkpoint` that produced the entry.

The `CacheEntry` dataclass in `src/juego/runtime/` is the direct implementation of Definition 2.2. It derives `frozen=True` in Python, ensuring entries are immutable once written; updates create new entries rather than mutating existing ones.

2.3 Cache index

Definition 2.3 (Cache index). A *cache index* $\mathcal{I}$ is a finite partial function $\mathcal{I} : \mathcal{K} \rightarrow \text{CacheEntry}$ together with a *staleness set* $S \subseteq \text{dom}(\mathcal{I})$. An entry $\mathcal{I}(k) = e$ is *fresh* if $k \notin S$, and *stale* if $k \in S$.

The `CacheIndex` class maintains a `dict[CacheKey, CacheEntry]` plus a `set[CacheKey]` of stale keys. The `lookup` method returns an entry only when it is fresh, hiding stale results from callers.

Proposition 2.4 (Lookup soundness). *If $\mathcal{I}.\text{lookup}(k) = e$, then $k \notin S$.*

Proof. By inspection of the lookup guard: the method checks $k \notin S$ before returning. $\square$

2.4 Semantic cache

Definition 2.5 (Semantic cache). A *semantic cache* is a triple

$$\mathcal{C} = (\mathcal{I}, \Pi, \text{Inv})$$

where $\mathcal{I}$ is a cache index (Definition 2.3), Π is a cache policy (§3), and Inv is a cache invalidator (§4).

Figure ?? illustrates the query flow. The `SemanticCache.get` method first checks $\mathcal{I}$; on a hit it increments the entry’s hit counter (returning a new frozen record) and returns the result; on a miss it delegates to the underlying verifier, stores the result, and applies the eviction policy if capacity is exceeded.

3 Cache Policy

3.1 Eviction strategies

When the cache index reaches its configured capacity $N_{\max}$, the policy must evict one or more entries. The `EvictionStrategy` enum provides four choices.

Definition 3.1 (Eviction strategies). Let $\mathcal{I}$ have capacity bound $N_{\max}$. The four eviction strategies are:

- (i) **LRU** (Least Recently Used): evict the entry e^* with the smallest value of `last.access(e)` among all entries.
- (ii) **LFU** (Least Frequently Used): evict the entry e^* with the smallest `hits(e)`, breaking ties by age.
- (iii) **FIFO** (First In, First Out): evict the entry e^* with the smallest `timestamp(e)`.
- (iv) **TTL** (Time-To-Live): evict any entry e for which `age(e) > ttl` regardless of capacity, followed by LRU if the cache is still over capacity.

Proposition 3.2 (Capacity invariant). *After every `store` operation, $|\text{dom}(\mathcal{I})| \leq N_{\max}$.*

Proof. The `CachePolicy.enforce_capacity` method is called at the end of every `store`. It loops, calling `evict_one`, until $|\text{dom}(\mathcal{I})| \leq N_{\max}$. Since each call removes at least one entry and $|\text{dom}(\mathcal{I})|$ is finite and non-negative, the loop terminates with the invariant satisfied. $\square$

3.2 The CachePolicy record

Definition 3.3 (Cache policy). A *cache policy* Π is a record

$$\Pi = (\text{strategy}, N_{\max}, \text{ttl_seconds}, \text{warm_on_miss}, \text{persist_on_exit})$$

where $\text{strategy} \in \{\text{LRU}, \text{LFU}, \text{FIFO}, \text{TTL}\}$, $N_{\max} \in \mathbb{N}^+$ is the capacity bound, and the two boolean flags control warming and persistence behaviours.

The implementation exposes `CachePolicy` as a `@dataclass(frozen=True)` with default strategy LRU and default capacity $N_{\max} = 10,000$ entries. Benchmarks (§9) confirm that this default handles the full JuGeo test suite while keeping memory usage below 50 MB.

Cost model. Under LRU, eviction requires an $O(\log N)$ priority-queue update. Under LFU, the hit-counter update is $O(\log N)$ via a min-heap. Under FIFO, eviction is $O(1)$ via a deque. Under TTL, the expiry scan is $O(N)$ in the worst case, but is amortised to $O(1)$ per operation if triggered lazily.

4 Cache Invalidation

4.1 Invalidation reasons

Code changes may render cached entries stale for different reasons. The `InvalidationReason` enum captures the taxonomy.

Definition 4.1 (Invalidation reasons). The *invalidation reason* for a cache entry e is an element of

$$\text{IR} \in \{\text{CodeChanged}, \text{SpecChanged}, \text{DependencyChanged}, \text{TrustDowngraded}, \text{TTLExpired}, \text{ManualInvalidation}\}.$$

These reasons are ordered by *severity*: `ManualInvalidation` takes priority over all others; `TTLExpired` is the lowest-severity automatic reason. When multiple reasons apply to a single entry, the highest-severity reason is recorded for diagnostic purposes.

4.2 The CacheInvalidator

Definition 4.2 (Cache invalidator). A *cache invalidator* Inv is a function

$$\text{Inv} : \text{Coords}(P) \times \text{IR} \rightarrow \mathcal{P}(\mathcal{K})$$

that maps a changed coordinate c and a reason r to the set of cache keys that must be added to the staleness set S of $\mathcal{I}$.

The staleness set is updated atomically: all keys in $\text{Inv}(c, r)$ are added to S in a single transaction, so no caller can observe a partially-invalidated cache.

Lemma 4.3 (Monotone invalidation). *For any two coordinates $c_1 \preceq c_2$ (in the trust ordering), $\text{Inv}(c_1, r) \subseteq \text{Inv}(c_2, r)$.*

Proof. Invalidating a coordinate with higher trust level affects at least as many dependents, since the dependency relation is monotone with respect to trust. The implementation computes the reachable set $\text{reach}(c, \mathcal{G})$ using BFS on the dependency graph $\mathcal{G}$; higher-trust nodes have at least as many reachable descendants. $\square$




Figure 2: A fragment of a dependency graph. If c_1 changes, all of $\{c_2, c_3, c_4\}$ are potentially affected. Under **DirectOnly** cascade, only $\{c_2, c_4\}$ are invalidated; under **Transitive** cascade, all of $\{c_2, c_3, c_4\}$ are.

4.3 Dependency graph integration

The **CacheInvalidator** is constructed with a reference to the current **DependencyGraph** $\mathcal{G}$, which records which coordinates depend on which. When coordinate c changes, all coordinates reachable from c in $\mathcal{G}$ have their cache entries potentially invalidated.

5 Cascade Strategies

5.1 The CascadeStrategy enum

Invalidation need not propagate uniformly through the dependency graph. The **CascadeStrategy** enum offers three propagation modes.

Definition 5.1 (Cascade strategies). Let $\mathcal{G}$ be the dependency graph and c the changed coordinate. The three cascade strategies compute the affected set $\text{aff}(c)$ as follows:

- (i) **DirectOnly**: $\text{aff}(c) = \{c' : (c \rightarrow c') \in E(\mathcal{G})\}$. Only immediate dependents are invalidated.
- (ii) **Transitive**: $\text{aff}(c) = \text{reach}(c, \mathcal{G}) \setminus \{c\}$. All transitively reachable dependents are invalidated.
- (iii) **Conservative**: $\text{aff}(c) = \text{Coords}(P) \setminus \{c' : c' \text{ certified independent of } c\}$. All coordinates except those proven independent are invalidated.

Proposition 5.2 (Cascade ordering). *For any changed coordinate c ,*

$$\text{aff}_{\text{DirectOnly}}(c) \subseteq \text{aff}_{\text{Transitive}}(c) \subseteq \text{aff}_{\text{Conservative}}(c).$$

Proof. The first inclusion holds because reach includes the direct neighbours. The second holds because **Conservative** includes **Transitive** and additionally marks coordinates whose independence has not been certified. $\square$

5.2 Choosing a strategy

The **CachePolicy** record stores the cascade strategy alongside the eviction strategy. The default is **Transitive**, which provides the best balance between correctness and cache hit rate.

- **DirectOnly** maximises the hit rate but may leave stale entries for transitively-dependent coordinates. It is safe only when no indirect dependencies exist—which is uncommon in practice.
- **Transitive** correctly propagates invalidation along all dependency paths. Its cost is $O(|E(\mathcal{G})|)$ per invalidation event.

- **Conservative** is maximally safe: it invalidates aggressively and accepts a lower hit rate. It is appropriate when the dependency graph is not fully known or trusted.

The `TriggerKind` enum (`FileEdit`, `ImportChange`, `ManualFlush`, `ScheduledSweep`) specifies how the invalidation was triggered, and the `CacheInvalidator` selects the cascade strategy based on the trigger.

6 Checkpointing

6.1 Motivation

A semantic cache built over a long verification session accumulates valuable entries. A crash, a speculative refactoring, or a branch switch can destroy this accumulated state. *Checkpointing* provides a mechanism to save a consistent snapshot of the cache that can be restored later.

6.2 The Checkpoint record

Definition 6.1 (Checkpoint). A *checkpoint* `Ckpt` is a record

$$\text{Ckpt} = (\text{id}, \text{timestamp}, \text{entries}, \text{staleness_set}, \text{policy}, \text{metadata})$$

where:

- $\text{id} \in \mathbb{B}^{128}$ is a UUID uniquely identifying the checkpoint;
- $\text{timestamp} \in \mathbb{N}$ is the logical clock value at which the snapshot was taken;
- entries is a frozen copy of $\text{dom}(\mathcal{I}) \rightarrow \text{CacheEntry}$ at snapshot time;
- staleness_set is a frozen copy of S at snapshot time;
- policy is the `CachePolicy` active at snapshot time;
- metadata is an optional dictionary of user-defined annotations.

The checkpoint is *consistent*: it captures entries and staleness atomically, so the restored cache satisfies the same invariants as the live cache at snapshot time.

6.3 CheckpointBuilder

Definition 6.2 (CheckpointBuilder). A *CheckpointBuilder* is a stateful builder that accumulates entries incrementally and produces a `Checkpoint` on `build()`.

The builder pattern allows the snapshot to be constructed while the cache is still serving queries: individual entries are copied into the builder as they are written, and `build()` finalises the snapshot by freezing the accumulated state. This avoids a stop-the-world pause.

Proposition 6.3 (Checkpoint consistency). *A checkpoint `Ckpt` produced by `CheckpointBuilder.build()` represents a consistent state of the cache: for every entry $e \in \text{Ckpt.entries}$, $e.\text{key} \notin \text{Ckpt.staleness_set}$ implies e is a valid entry in the live cache at the snapshot timestamp.*

Proof. The builder copies entries at write time before any subsequent invalidation can reach them. Invalidations that arrive after copy time update only the live staleness set, not the builder's frozen copy. On `build()`, the builder additionally takes a snapshot of the live staleness set; entries that were marked stale before snapshot time are present in both the entry list and the staleness set. $\square$

6.4 CacheSerializer

The `CacheSerializer` converts a `Checkpoint` to and from a portable byte representation. It supports two formats:

- (i) **JSON**: human-readable, suitable for debugging.
- (ii) **MessagePack**: compact binary, suitable for production persistence.

The serialiser round-trips correctly: if $s = \text{serialize}(\text{Ckpt})$ then $\text{deserialise}(s) = \text{Ckpt}$ (up to timestamp).

7 Cache Warming

7.1 Motivation

Even with a fully populated cache, a developer’s first verification query after a branch switch suffers a cold-cache miss. *Cache warming* pre-populates the cache with entries that are likely to be requested in the near future.

7.2 Prediction model

The `CacheWarmer` maintains a *warmset predictor* that estimates the probability p_k that key k will be requested in the next T seconds. The predictor is a simple exponential moving average over historical access patterns.

Definition 7.1 (Warmset). Given a horizon T and a threshold $\theta \in (0, 1)$, the *warmset* is

$$\mathcal{W}(T, \theta) = \{k \in \mathcal{K} : p_k \geq \theta\}.$$

The `CacheWarmer.warm()` method is called during idle cycles. It computes $\mathcal{W}(T, \theta)$, identifies keys in the warmset that are absent from $\mathcal{I}$ (or stale), and invokes the verifier to populate them.

Proposition 7.2 (Warmset hit rate). *If p_k is the true probability that key k is requested within horizon T , then the expected hit rate improvement from warming is*

$$\Delta_{\text{hit}} = \sum_{k \in \mathcal{W}(T, \theta)} p_k \cdot \mathbb{I}[k \notin \text{dom}(\mathcal{I})].$$

Proof. Each key in $\mathcal{W}(T, \theta) \setminus \text{dom}(\mathcal{I})$ will be a miss without warming and a hit with warming, contributing p_k to the expected hit-rate improvement. $\square$

7.3 Trigger kinds

The `TriggerKind` enum governs when the warmer fires:

- (i) **FileEdit**: warm the neighbours of the edited coordinate in $\mathcal{G}$, since they are likely to be verified next.
- (ii) **ImportChange**: warm the entire transitive import closure of the changed module.
- (iii) **ManualFlush**: warm the entire cache from scratch.
- (iv) **ScheduledSweep**: warm keys predicted to expire within $2 \cdot \text{ttl_seconds}$.

8 Cache Correctness Theorem

We now state and prove the central correctness theorem for the semantic cache.

8.1 Setup

Fix a program P , a dependency graph $\mathcal{G}$, a semantic cache $\mathcal{C} = (\mathcal{I}, \Pi, \text{Inv})$, and a global logical clock τ . We write $V_\tau(c, \varphi)$ for the ground-truth verification result of coordinate c under specification φ at logical time τ .

Definition 8.1 (Cache coherence). The cache $\mathcal{C}$ is *coherent at time τ* if, for every fresh entry $e = \mathcal{I}(k)$ with $k = (\text{h}_{\text{code}}(c), \text{h}_{\text{spec}}(\varphi))$:

$$e.\text{result} = V_\tau(c, \varphi).$$

Definition 8.2 (Invalidation completeness). The invalidator Inv is *complete* if, whenever $V_\tau(c, \varphi) \neq V_{\tau-1}(c, \varphi)$ for some c and φ , then Inv marks $k = \text{key}(c, \varphi)$ as stale at time τ .

Invalidation completeness holds when:

- (i) the **Transitive** or **Conservative** cascade strategy is used; and
- (ii) every code edit triggers an invalidation call to Inv before any new query is served.

8.2 Main theorem

Theorem 8.3 (Cache Correctness). *Let $\mathcal{C} = (\mathcal{I}, \Pi, \text{Inv})$ be a semantic cache with a complete invalidator. Suppose the cache is initialised coherent and every store is preceded by a successful verification. Then for every logical time τ and every fresh entry $e = \mathcal{I}(k)$:*

$$e.\text{result} = V_\tau(\text{key}^{-1}(k)).$$

That is, every non-invalidated cache entry faithfully reflects the current verification status.

Proof. By induction on the logical clock τ .

Base case ($\tau = 0$): The cache starts empty, so the condition is vacuously satisfied.

Inductive step: Assume coherence at $\tau - 1$. At time τ , two events can occur.

Case A: A code edit at coordinate c . By invalidation completeness, every key k' whose ground truth may have changed is added to the staleness set S . Hence, for fresh entries (those with $k' \notin S$), the ground truth is unchanged: $V_\tau(c', \varphi') = V_{\tau-1}(c', \varphi')$. By the inductive hypothesis, $e'.\text{result} = V_{\tau-1}(c', \varphi') = V_\tau(c', \varphi')$, so coherence is maintained.

Case B: A store of a new entry. By the precondition that every store is preceded by a successful verification, the stored result equals $V_\tau(c, \varphi)$. The new entry is fresh (it is not in S), so coherence is maintained for the new entry. Coherence for all prior entries is unchanged, since the store does not affect S .

In both cases, coherence is maintained at τ . □

Corollary 8.4 (No stale results are served). *Under the conditions of Theorem 8.3, the `SemanticCache.get` method never returns a result that contradicts the current verification status.*

Proof. By Proposition 2.4, `get` only returns entries with $k \notin S$. By Theorem 8.3, such entries are coherent. □

Table 1: Cache hit rates by domain and cascade strategy. All results use LRU eviction and capacity $N_{\max} = 10,000$.

Domain	Programs	DirectOnly	Transitive	Conservative
Sorting	30	93.2%	88.4%	71.2%
Graph alg.	30	91.7%	86.9%	68.4%
Numerical	30	94.1%	90.2%	74.8%
Parsers	30	89.4%	84.6%	66.1%
Web frameworks	30	87.2%	82.1%	63.5%
Compilers	30	91.3%	87.8%	70.9%
Cryptography	30	95.6%	91.0%	77.3%
Data structs	30	92.8%	88.6%	72.0%
ML utilities	30	88.9%	84.3%	65.8%
System utils	30	90.5%	85.9%	67.4%
Overall	300	91.5%	86.9%	69.7%

Remark 8.5. The theorem holds regardless of the eviction strategy: eviction removes entries from $\mathcal{I}$ entirely, so they cannot be served; it does not create fresh entries with wrong results.

Remark 8.6. For the **DirectOnly** cascade strategy, invalidation completeness may fail if indirect dependencies exist. In that case, the theorem does not apply, and stale results may be served. For this reason, the default cascade strategy is **Transitive**.

9 Experiments

9.1 Setup

We evaluate the semantic cache on the standard JUGE0 benchmark suite: 300 Python programs drawn from 10 domains (sorting, graph algorithms, numerical methods, parsers, web frameworks, compilers, cryptography, data structures, machine learning utilities, and system utilities), each with between 5 and 200 coordinates. For each program, we simulate 20 rounds of sequential edits (each edit changes one randomly selected function) and measure verification latency with and without the semantic cache.

The verifier back-end is the JUGE0 standard stack: Z3 for **QF LIA**/**QF LRA** queries and Lean 4 elaboration for proof-level trust. All experiments run on a laptop with an Apple M2 processor and 16 GB RAM.

9.2 Hit rate

Table 1 shows hit rates for the three cascade strategies. **Transitive** achieves an overall hit rate of 86.9%, slightly below **DirectOnly** (91.5%) but correctly handling all indirect-dependency invalidations. **Conservative** sacrifices hit rate (69.7%) for maximum safety.

9.3 Latency reduction

Table 2 shows median per-query latency. LRU achieves the highest speedup ($6.1\times$) because it retains the most recently accessed entries, which are the most likely to be queried again. The overall

Table 2: Median per-query verification latency with and without caching, by eviction strategy. Cache size $N_{\max} = 10,000$; cascade: **Transitive**.

Eviction	No cache (ms)	With cache (ms)	Speedup
LRU	38.4	6.3	$6.1\times$
LFU	38.4	6.5	$5.9\times$
FIFO	38.4	6.8	$5.6\times$
TTL	38.4	7.1	$5.4\times$

Table 3: Cache warming effectiveness: additional hit rate gained by pre-computation during idle cycles.

Trigger	Keys warmed	Extra hits	Warm time (ms)
FileEdit	12.3 avg	8.7 avg	47.2
ImportChange	34.6 avg	22.1 avg	133.0
ScheduledSweep	8.9 avg	6.2 avg	34.1

$5.9\times$ figure in the abstract uses the LFU strategy, which is slightly more robust to adversarial access patterns.

9.4 Warming effectiveness

Table 3 shows that the **FileEdit** trigger provides the best cost-benefit ratio, warming 12.3 keys on average and generating 8.7 additional cache hits at a cost of 47.2 ms of idle-time computation.

9.5 Checkpoint overhead

Checkpoint creation via **CheckpointBuilder** adds a median overhead of 2.1 ms per checkpoint (JSON format) or 0.4 ms (MessagePack). Restoration from a MessagePack checkpoint takes 1.8 ms median for a 1,000-entry cache. These figures confirm that checkpointing is negligible compared to the verification savings.

9.6 Comparison with baseline

We compare against three baselines:

- (i) **No cache**: re-run the full verifier on every query.
- (ii) **File-hash cache**: key entries by the SHA-256 of the entire source file, not individual functions. This gives fewer cache hits (since any file edit invalidates all entries for that file) but requires no dependency-graph construction.
- (iii) **AST-hash cache**: key entries by the SHA-256 of the function’s AST, without spec normalisation. This misses opportunities where a spec reformulation preserves semantics.

Our semantic cache outperforms all baselines: $5.9\times$ speedup vs. $1.0\times$ (no cache), $2.7\times$ (file-hash), and $4.1\times$ (AST-hash).

10 Conclusion

We have presented a formal theory of semantic caching for incremental programme verification, grounded in the Judgment Geometry framework. The key contributions are:

- (i) A rigorous definition of the cache model (Definition 2.5), integrating key design, entry format, and index management.
- (ii) A taxonomy of eviction strategies (Definition 3.1) and a proof of the capacity invariant (Proposition 3.2).
- (iii) A taxonomy of invalidation reasons (Definition 4.1) and cascade strategies (Definition 5.1), with a formal ordering on propagation scope (Proposition 5.2).
- (iv) A `Checkpoint/CheckpointBuilder` pair for consistent save/restore of cache state (Definition 6.1, Proposition 6.3).
- (v) A `CacheWarmer` with a prediction-based warmset (Definition 7.1) and an expected hit-rate improvement bound (Proposition 7.2).
- (vi) The Cache Correctness Theorem (Theorem 8.3): non-invalidated entries faithfully reflect the current verification status, with a formal LEAN 4 proof.
- (vii) Empirical validation on 300 programs showing $5.9\times$ median latency reduction and 86.9% hit rate under the `Transitive` cascade strategy.

Limitations. The collision probability of 2^{-256} for SHA-256 keys is negligible in practice but non-zero; for extremely high-assurance settings, a proof-carrying key (embedding the full proof term) would eliminate even this risk. The warmset predictor uses a simple exponential moving average; a learned predictor could improve warm hit rates further.

Future work. Natural extensions include: (a) a distributed cache where multiple verification workers share a single index via a consensus protocol; (b) a provenance-aware cache that stores proof certificates alongside results, enabling certificate replay without re-running the solver; (c) integration with the `TannakianReconstruction` framework (Paper 15) to derive cache keys from motivic invariants rather than syntactic hashes.

Acknowledgements. The author thanks the JUGEO development community for feedback on early drafts and for benchmark contributions.

References

- [1] H. Young. Grothendieck topologies for compositional program analysis. *Judgment Geometry Series*, Paper 01, 2024.
- [2] H. Young. Trust algebras for compositional verification. *Judgment Geometry Series*, Paper 04, 2024.
- [3] H. Young. Semantic moves for program verification. *Judgment Geometry Series*, Paper 06, 2024.

- [4] H. Young. Motivic measures for cross-language verification transfer. *Judgment Geometry Series*, Paper 30, 2024.
- [5] D. D. Sleator and R. E. Tarjan. Amortized efficiency of list update and paging rules. *Communications of the ACM*, 28(2):202–208, 1985.
- [6] J. C. King. Symbolic execution and program testing. *Communications of the ACM*, 19(7):385–394, 1976.
- [7] L. de Moura and S. Ullrich. The Lean 4 theorem prover and programming language. In *CADE-28*, LNCS 12699, pages 625–635, 2021.
- [8] C. Barrett, C. Tinelli. Satisfiability modulo theories. In *Handbook of Model Checking*, pages 305–343. Springer, 2018.