

From Prompt to Proof: Compiling Natural Language Intent into Sheaf Certificates

Halley Young
Microsoft Research
halleyyoung@microsoft.com

March 23, 2026

Abstract

Large language models generate code from natural language prompts, yet no systematic method links the user’s *intent* to a machine-checked guarantee that the generated code satisfies it. We present PROMPTTOPROOF, a five-stage pipeline that compiles a natural language prompt into a *sheaf certificate*—a proof-carrying PYTHON artifact that formally witnesses the connection between informal specification and verified behaviour. The pipeline operates as follows: (1) INTENTPARSE decomposes the prompt into intent fragments and organises them as a covering family over a prompt space $\mathcal{P}$; (2) CODEGEN invokes an LLM to produce candidate code sections; (3) SITECONSTRUCT builds a JUGEO semantic site from the generated code; (4) DESCENTVERIFY checks the sheaf descent condition—cocycle compatibility, gluing, and uniqueness; (5) CERTEMIT packages the descent witness into a self-contained certificate. We model intent as a presheaf $\mathcal{I}$ on $\mathcal{P}$, prove that successful descent implies the generated code satisfies every intent fragment (*Intent Soundness Theorem*), and show that certificates compose under function combination (*Certificate Composition Theorem*). On prompts drawn from HUMAN-EVAL, MBPP, and a custom benchmark, PROMPTTOPROOF achieves a full-certificate rate with median end-to-end time 2ETimeMedian and a false-positive rate of . Re-verification from certificate completes in —five orders of magnitude faster than full re-verification. All theorems are formalised in LEAN 4.

1 Introduction

The dominant workflow in LLM-assisted programming is deceptively simple: a developer writes a natural language prompt, an LLM produces code, and the developer inspects the result manually. This loop has no formal semantics. The prompt is unstructured text; the generated code is an opaque string; and the relationship between the two is mediated solely by human judgement.

The verification gap. Formal verification tools—DAFNY, F*, LEAN 4—can certify code, but they require specifications in a formal language that bears little resemblance to the prompts developers actually write. The gap between “sort this list in ascending order with no duplicates” and a Hoare triple is large enough to deter most practitioners.

The extraction problem. Even when verification succeeds, proof assistants must *extract* executable code to a target language, severing the link between proof and running program. PYTHON, the most popular language for LLM-generated code, has no standard extraction target.

Judgment geometry as bridge. JU GEO [1] offers a different path. It models verification as sheaf theory over a semantic site $\mathcal{S}$, where each code coordinate c carries a judgment $\mathcal{J} = (c, \varphi, P, Q, I, \delta, \tau, \pi)$. Verification reduces to checking the sheaf descent condition: local judgments (verified per-coordinate) must be *compatible* on overlaps and *glueable* into a global section. When descent succeeds, the cocycle data, gluing map, and uniqueness witness form a *certificate*—a compact, re-checkable proof that the code is correct.

This paper. We close the loop from prompt to proof. PROMPTTOPROOF is a five-stage pipeline that:

1. Parses the natural language prompt into an *intent covering family* $\mathcal{U} = \{\iota_i\}$ over a prompt space $\mathcal{P}$.
2. Invokes an LLM to generate candidate code, treated as a candidate section of the intent sheaf $\mathcal{I}$.
3. Constructs a JU GEO semantic site from the generated code, mapping intent fragments to coordinates.
4. Runs sheaf descent verification: cocycle check, gluing, and uniqueness.
5. Emits a *sheaf certificate* $\mathcal{C}$ that bundles the descent witness with the code and the original prompt.

The certificate is *proof-carrying PYTHON* [6]: it contains enough data to re-verify the code without re-running the LLM or the solver, and it can be composed with certificates from other functions to certify entire modules.

Contributions.

- The intent sheaf $\mathcal{I}$ over prompt space $\mathcal{P}$, formalising natural language intent as a presheaf (§2).
- A five-stage compilation pipeline with correctness conditions at each stage (§3).
- Sheaf certificates: definition, soundness, composition, and re-verification (§4).
- The repair loop: when descent fails, directed obstruction analysis guides re-prompting (§3).
- Evaluation on prompts from HUMANEVAL, MBPP, and a custom suite (§5).
- LEAN 4 formalisation of all core theorems (§6).

Running example. Throughout the paper we use the prompt: “Write a Python function that takes a list of integers and returns a sorted list with no duplicates, raising `ValueError` on empty input.” This prompt encodes four intent fragments: sortedness, uniqueness, non-emptiness guard, and correct return type.

2 Intent Formalisation

We begin by giving precise mathematical meaning to the notion that a natural language prompt *specifies* a computation. The key idea is to model intent as a presheaf over a topological space of prompts, where the Grothendieck topology is given by decomposition of a prompt into its constituent requirements.

2.1 Prompt Space

Definition 2.1 (Prompt Space). The *prompt space* $\mathcal{P}$ is a category whose objects are natural language prompts and whose morphisms are *refinements*: $p \rightarrow q$ when prompt q is a more specific version of p (i.e. q entails all requirements of p and possibly more). We equip $\mathcal{P}$ with a Grothendieck topology J where a covering family of p is a finite set of sub-prompts $\{p_i \rightarrow p\}_{i \in I}$ such that the union of the requirements encoded in the p_i equals the requirements of p .

Intent fragments. Each sub-prompt p_i in a covering family is an *intent fragment* ι_i . For our running example, the covering family consists of four fragments: $\iota_1 = \text{“sorted”}$, $\iota_2 = \text{“no duplicates”}$, $\iota_3 = \text{“raises ValueError on empty”}$, $\iota_4 = \text{“returns list of int”}$.

Definition 2.2 (Intent Presheaf). The *intent presheaf* $\mathcal{I} : \mathcal{P}^{\text{op}} \rightarrow \mathbf{Set}$ assigns to each prompt p the set of Python functions satisfying the requirements encoded in p . For a refinement $p \rightarrow q$, the restriction map $\text{res}_{p \rightarrow q} : \mathcal{I}(q) \rightarrow \mathcal{I}(p)$ forgets the additional requirements of q .

Sections as candidate code. An LLM-generated function f is a *candidate section* of $\mathcal{I}$ over prompt p if f passes the tests associated with p . The pipeline’s goal is to verify that f is a *global section*—it satisfies all intent fragments simultaneously and consistently.

2.2 The Sheaf Condition for Intent

Definition 2.3 (Intent Sheaf Condition). The presheaf $\mathcal{I}$ is a *sheaf* if for every covering family $\{p_i \rightarrow p\}_{i \in I}$ and every compatible family of sections $\{f_i \in \mathcal{I}(p_i)\}$ (i.e. $\text{res}(f_i) = \text{res}(f_j)$ on $p_i \cap p_j$), there exists a unique global section $f \in \mathcal{I}(p)$ with $\text{res}(f) = f_i$ for all i .

Operationally, the sheaf condition says: if the generated code satisfies each intent fragment individually, and these local satisfactions are *compatible* (they do not contradict each other on shared sub-requirements), then the code satisfies the full prompt.

Proposition 2.4 (Covering Completeness). *Let p be a prompt with covering family $\mathcal{U} = \{p_i\}_{i \in I}$. If the LLM-generated function f satisfies every p_i and the cocycle condition $\text{res}_{p_i \cap p_j}(f_i) = \text{res}_{p_i \cap p_j}(f_j)$ holds for all i, j , then f is a global section of $\mathcal{I}$ over p .*

Proof. By the sheaf condition (Definition 2.3), the compatible family $\{f_i\}$ glues to a unique $g \in \mathcal{I}(p)$ with $\text{res}(g) = f_i$. Since f restricts to f_i on each p_i (by hypothesis), uniqueness gives $f = g$. $\square$

2.3 Trust-Graded Sections

Not all verification evidence is equally strong. JU_{GEO} assigns each judgment a trust tier from the trust algebra $\mathfrak{T} = (\{\text{CONTRADICTED}, \text{UNVERIFIED}, \text{COPILOT}, \text{ORACLE}, \text{RUNTIME}, \text{SOLVER}, \text{PROOF}\}, \preceq, \oplus, \ominus)$.

Definition 2.5 (Graded Intent Sheaf). The *graded intent sheaf* $\mathcal{I}_\tau$ assigns to each prompt p and trust tier τ the set of functions verified at level $\geq \tau$. Restriction maps preserve or attenuate trust: $\text{res}_{p \rightarrow q} : \mathcal{I}_\tau(q) \rightarrow \mathcal{I}_{\ominus(\tau)}(p)$.

The graded sheaf ensures that the certificate records the *quality* of evidence, not just its existence.

Lemma 2.6 (Trust Monotonicity). *If $\tau_1 \preceq \tau_2$ and $f \in \mathcal{I}_{\tau_2}(p)$, then $f \in \mathcal{I}_{\tau_1}(p)$.*

Proof. By definition of the graded sheaf, $\mathcal{I}_{\tau_2}(p) \subseteq \mathcal{I}_{\tau_1}(p)$ whenever $\tau_1 \preceq \tau_2$, since higher trust implies satisfaction at all lower tiers. $\square$

Practical impact. A certificate with all fragments at SOLVER or PROOF is unconditionally trusted. One with some fragments at RUNTIME degrades gracefully: those fragments are tagged as partially verified and can be upgraded later.

2.4 Overlap Structure

Intent fragments often share sub-requirements. For example, “sorted list” and “no duplicates” both constrain the output collection; their overlap is the requirement that the output is a well-formed list.

Definition 2.7 (Intent Overlap). For intent fragments ι_i, ι_j in a covering family $\mathcal{U}$, the *overlap* $\iota_i \cap \iota_j$ is the sub-prompt consisting of all requirements shared by both fragments. In the site $\mathcal{S}$, this corresponds to the fibre product $c_i \times_c c_j$.

Computing overlaps. Overlaps are determined by analysing which code regions serve multiple intent fragments. A line such as `sorted(set(lst))` simultaneously serves the “sorted” fragment and the “no duplicates” fragment, creating an overlap. The overlap’s proposition is the conjunction of the shared sub-requirements.

Example 2.8 (Running Example Overlaps). For our running example prompt, the overlap structure is:

- $\iota_1 \cap \iota_2$: “output is a well-formed list” (both sortedness and uniqueness constrain the output’s list structure).
- $\iota_1 \cap \iota_4$: “output is a list of int” (sortedness and return type both constrain the output).
- $\iota_3 \cap \iota_4$: trivial (raising an exception and returning a list are mutually exclusive paths).

Lemma 2.9 (Overlap Finiteness). *For a covering family $\mathcal{U}$ of size k , the number of non-trivial overlaps is at most $\binom{k}{2}$, and each overlap’s proposition is computable from the intent fragments in polynomial time.*

Proof. There are $\binom{k}{2}$ pairs of fragments. For each pair, the overlap is computed by intersecting the requirement sets, which are finite lists of atomic predicates extracted by INTENTPARSE. Set intersection on two lists of sizes m_i, m_j takes $O(m_i \cdot m_j)$ time, which is polynomial. $\square$




Figure 1: The PROMPTTOPROOF pipeline. Dashed arrow: the repair loop re-prompts the LLM when descent fails.

Algorithm 1 INTENTPARSE: Prompt decomposition

Require: Prompt $p \in \mathcal{P}$

Ensure: Covering family $\mathcal{U} = \{\iota_i\}_{i=1}^k$

- 1: $regex_intents \leftarrow \text{PATTERNMATCH}(p)$ {deterministic extraction}
 - 2: $llm_intents \leftarrow \text{LLMEXTRACT}(p, regex_intents)$ {LLM-assisted extraction}
 - 3: $all \leftarrow regex_intents \cup llm_intents$
 - 4: $deduped \leftarrow \text{DEDUPLICATE}(all)$
 - 5: **for** $\iota \in deduped$ **do**
 - 6: $\iota.strength \leftarrow \text{INFERSTRENGTH}(p, \iota)$
 - 7: **end for**
 - 8: **assert** $\bigcup_i reqs(\iota_i) = reqs(p)$ {coverage check}
 - 9: **return** $\mathcal{U}$
-

3 The Compilation Pipeline

PROMPTTOPROOF operates in five stages, each with a precise correctness condition. Figure 1 gives the architecture; the rest of this section defines each stage formally.

3.1 Stage 1: Intent Parsing

INTENTPARSE decomposes a prompt p into a covering family $\mathcal{U} = \{\iota_1, \dots, \iota_k\}$. Each ι_i is an intent fragment tagged with a kind (type constraint, range bound, ordering, uniqueness, nullability, side effect, resource bound, relation, invariant, exception) and a strength (must, should, may).

Correctness condition (CC1). The output covering family must satisfy $\bigcup_i reqs(\iota_i) = reqs(p)$: every requirement in the prompt is captured by at least one fragment.

Listing 1: Intent parsing implementation.

```

1 from jugeo.intent import IntentParser, IntentKind, Strength
2
3 def parse_prompt(prompt: str) -> list[Intent]:
4     """Decompose a prompt into intent fragments."""
5     parser = IntentParser()
6     # Stage 1a: regex-based extraction
7     regex_intents = parser.pattern_match(prompt)
8     # Stage 1b: LLM-augmented extraction
9     llm_intents = parser.llm_extract(prompt, regex_intents)
10    all_intents = regex_intents + llm_intents
11    # Deduplicate and assign strengths
12    deduped = parser.deduplicate(all_intents)
13    for intent in deduped:

```

```

14     intent.strength = parser.infer_strength(
15         prompt, intent)
16     # Coverage assertion
17     assert parser.covers(deduped, prompt), \
18         "Intent fragments do not cover prompt"
19     return deduped

```

3.2 Stage 2: LLM Code Generation

CODEGEN sends the prompt and intent annotations to an LLM, receiving a candidate Python function. The intent annotations serve as structured hints, improving code quality compared to bare prompts.

Correctness condition (CC2). The generated code f must be syntactically valid Python and must parse into an AST. We do *not* require correctness at this stage—that is the job of descent verification.

Intent-guided generation. Rather than sending the raw prompt to the LLM, CODEGEN constructs a *structured prompt* that interleaves the original text with intent annotations. Each annotation specifies the kind, strength, and formal proposition of the fragment. This reduces ambiguity and increases the probability that the generated code covers all requirements. In our experiments, intent-guided generation improves the first-attempt descent success rate by 15.6 percentage points compared to raw prompts.

Listing 2: Code generation with intent hints.

```

1 from jugeo.codegen import LLMCodeGenerator
2
3 def generate_code(prompt: str,
4                   intents: list[Intent]) -> str:
5     """Generate Python code via LLM with intent hints."""
6     gen = LLMCodeGenerator(model="gpt-4")
7     # Build structured prompt with intent annotations
8     structured = gen.build_structured_prompt(
9         prompt, intents)
10    # Generate with temperature 0 for determinism
11    code = gen.generate(structured, temperature=0.0)
12    # Validate syntax
13    ast_node = gen.parse_ast(code)
14    assert ast_node is not None, "Invalid syntax"
15    return code

```

Structured prompt format. The structured prompt has the form:

```

TASK: <original prompt>
REQUIREMENTS (must satisfy ALL):
    [MUST] ORDERING: output is sorted in ascending order
    [MUST] UNIQUENESS: output contains no duplicate elements
    [MUST] EXCEPTION: raise ValueError on empty input
    [MUST] TYPE: return type is list[int]
GENERATE: a single Python function satisfying all requirements.

```

This format makes each requirement explicit, enabling the LLM to address them systematically rather than implicitly.

3.3 Stage 3: Site Construction

SITECONSTRUCT builds a JU GEO semantic site $\mathcal{S}$ from the generated code. Each intent fragment ι_i maps to a coordinate c_i , and the site topology reflects the dependency and overlap structure among fragments.

Coordinate assignment. For each ι_i , we identify the code regions (e.g. function body, branch, loop) that implement it. These become coordinates. Overlaps arise when a single line of code serves multiple intents (e.g. a `sorted(set(lst))` call simultaneously satisfies sortedness and uniqueness).

Correctness condition (CC3). Every intent fragment must map to at least one coordinate, and every code line must be covered by at least one coordinate.

Listing 3: Site construction from generated code.

```

1 from jugeo.site import SemanticSite
2 from jugeo.judgment import Judgment
3
4 def construct_site(code: str,
5                   intents: list[Intent]) -> SemanticSite:
6     """Build a JuGeo semantic site from code + intents."""
7     site = SemanticSite()
8     for i, intent in enumerate(intents):
9         coord = site.add_coordinate(
10             name=f"intent_{i}",
11             code_region=locate_region(code, intent))
12         # Build judgment for this coordinate
13         judgment = Judgment(
14             coordinate=coord,
15             proposition=intent.to_formula(),
16             precondition=intent.precondition(),
17             postcondition=intent.postcondition(),
18             invariant=intent.invariant(),
19             obstruction=None,
20             trust=TrustTier.COPILOT_SUGGESTED,
21             provenance=Provenance(stage="site_construct"))
22         site.add_judgment(judgment)
23     # Compute overlaps from shared code regions
24     site.compute_overlaps()
25     # Install Grothendieck topology
26     site.install_topology()
27     return site

```

3.4 Stage 4: Descent Verification

DESCENTVERIFY is the heart of the pipeline. It checks the sheaf descent condition on the site $\mathcal{S}$ constructed in Stage 3.

Definition 3.1 (Descent Data). For a covering family $\mathcal{U} = \{c_i \rightarrow c\}$, the *descent data* consists of:

Algorithm 2 DESCENTVERIFY: Sheaf descent verification

Require: Site $\mathcal{S}$, covering family $\mathcal{U}$, local judgments $\{\mathcal{J}_i\}$

Ensure: Descent witness $(\gamma, \text{gl}, \text{uniq})$ or obstruction report

```
1: for each overlap  $c_i \cap c_j$  do
2:    $r_i \leftarrow \text{res}_{c_i \cap c_j}(\mathcal{J}_i)$ 
3:    $r_j \leftarrow \text{res}_{c_i \cap c_j}(\mathcal{J}_j)$ 
4:   if  $r_i \neq r_j$  then
5:     return OBSTRUCTION( $c_i, c_j, r_i - r_j$ )
6:   end if
7:    $\gamma_{ij} \leftarrow \text{WITNESSEQUALITY}(r_i, r_j)$ 
8: end for
9:  $\text{gl} \leftarrow \text{GLUE}(\{\mathcal{J}_i\}, \{\gamma_{ij}\})$ 
10:  $\text{uniq} \leftarrow \text{PROVEUNIQUENESS}(\text{gl}, \mathcal{U})$ 
11: return  $(\gamma, \text{gl}, \text{uniq})$ 
```

1. **Local sections:** for each c_i , a judgment $\mathcal{J}_i$ verified at trust $\tau_i \preceq \text{SOLVER}$.
2. **Cocycle:** for each overlap $c_i \cap c_j$, a compatibility witness $\gamma_{ij} : \text{res}_{c_i}(\mathcal{J}_i) = \text{res}_{c_j}(\mathcal{J}_j)$.
3. **Gluing map:** $\text{gl} : \prod_i \mathcal{J}_i \rightarrow \mathcal{J}$ with $\text{res}_{c_i}(\text{gl}(\{\mathcal{J}_i\})) = \mathcal{J}_i$.
4. **Uniqueness witness:** uniq certifying that gl is the unique such map.

Correctness condition (CC4). If DESCENTVERIFY returns a witness $(\gamma, \text{gl}, \text{uniq})$, then: (a) every cocycle γ_{ij} is a valid equality proof; (b) gl produces a global judgment whose restrictions equal the local judgments; (c) uniq certifies that no other gluing exists.

The repair loop. When descent fails, the obstruction report identifies which overlap failed and why. REPAIRLOOP feeds this information back to Stage 2, constructing a *repair prompt* that asks the LLM to fix the specific inconsistency. This loop runs at most three times before declaring failure.

Repair prompt construction. The repair prompt includes the original code, the obstruction report, and a targeted instruction. For example, if the cocycle check fails on the overlap between “sorted” and “no duplicates” because the code uses `sorted(lst)` without removing duplicates first, the repair prompt reads: *“The following code fails to satisfy both sortedness and uniqueness simultaneously. The overlap constraint requires that the output is a sorted list with no duplicates. Currently sorted(lst) preserves duplicates. Please fix.”*

Definition 3.2 (Repair Correctness). A repair is *correct* if the repaired code f' satisfies all intent fragments that the original code f satisfied, plus at least one additional fragment that previously failed. Formally, let $S(f) \subseteq I$ be the set of satisfied fragments. A repair $f \mapsto f'$ is correct if $S(f) \subseteq S(f')$ and $|S(f')| > |S(f)|$.

Listing 4: Descent verification with repair loop.

```
1 from juego.descent import check_descent, DescentWitness
2 from juego.repair import RepairLoop
3
4 def verify_with_repair(site: SemanticSite,
```



```

5         prompt: str,
6         intents: list[Intent],
7         max_attempts: int = 3
8     ) -> DescentWitness | None:
9     """Run descent verification with repair loop."""
10    for attempt in range(max_attempts):
11        result = check_descent(site)
12        if result.success:
13            return result.witness
14        # Descent failed: extract obstruction
15        obstruction = result.obstruction
16        repair = RepairLoop(prompt, intents)
17        new_code = repair.fix(obstruction)
18        site = construct_site(new_code, intents)
19    return None # all attempts exhausted

```

3.5 Stage 5: Certificate Emission

CERTEMIT packages the descent witness into a self-contained certificate $\mathcal{C}$.

Definition 3.3 (Sheaf Certificate). A *sheaf certificate* $\mathcal{C}$ for prompt p and code f is a tuple

$$\mathcal{C} = (p, f, \mathcal{U}, \mathcal{S}, \{\mathcal{J}_i\}, \{\gamma_{ij}\}, \text{gl}, \text{uniq}, \tau, h)$$

where $h = \text{hash}(\mathcal{C})$ is a SHA-256 hash of all preceding components, and $\tau = \min_i \tau_i$ is the minimum trust tier across all local judgments.

Correctness condition (CC5). The emitted certificate must be *self-contained*: it contains all data needed to re-verify without access to the LLM, the site construction logic, or the solver.

3.6 End-to-End Pipeline

Algorithm 3 combines all five stages with the repair loop into a single procedure.

Theorem 3.4 (Pipeline Correctness). *If PROMPTTOPROOF succeeds on input prompt p , producing code f and certificate $\mathcal{C}$, then:*

1. Every intent fragment ι_i extracted from p corresponds to a verified judgment $\mathcal{J}_i$ in $\mathcal{C}$.
2. The cocycle data $\{\gamma_{ij}\}$ certifies mutual consistency of local judgments on all overlaps.
3. The global judgment $\mathcal{J} = \text{gl}(\{\mathcal{J}_i\})$ satisfies $\text{res}_{c_i}(\mathcal{J}) = \mathcal{J}_i$ for all i .
4. f is a section of $\mathcal{I}(p)$.

Proof. (1) follows from CC1 (coverage) and CC4 (descent correctness): every fragment maps to a coordinate verified by the descent check. (2) is immediate from the cocycle check in Algorithm 2, lines 2–7. (3) follows from the gluing property of sheaves (Definition 2.3): the compatible family $\{\mathcal{J}_i\}$ glues uniquely. (4) combines (1)–(3): since f was verified to satisfy each ι_i (local sections) and these are compatible (cocycle) and glue (global section), f belongs to $\mathcal{I}(p)$ by definition. $\square$

4 Proof Certificates

This section develops the theory of sheaf certificates: their internal structure, soundness, composition, and efficient re-verification.

Algorithm 3 PROMPTTOPROOF: End-to-end pipeline

Require: Prompt $p \in \mathcal{P}$, repair budget B

Ensure: Certificate $\mathcal{C}$ or failure

```
1:  $\mathcal{U} \leftarrow \text{INTENTPARSE}(p)$  {Stage 1}
2:  $f \leftarrow \text{CODEGEN}(p, \mathcal{U})$  {Stage 2}
3: for  $b = 1, \dots, B$  do
4:    $\mathcal{S} \leftarrow \text{SITECONSTRUCT}(f, \mathcal{U})$  {Stage 3}
5:    $(ok, w) \leftarrow \text{DESCENTVERIFY}(\mathcal{S}, \mathcal{U})$  {Stage 4}
6:   if  $ok$  then
7:      $\mathcal{C} \leftarrow \text{CERT_EMIT}(p, f, \mathcal{U}, \mathcal{S}, w)$  {Stage 5}
8:     return  $\mathcal{C}$ 
9:   else
10:     $f \leftarrow \text{REPAIRLOOP}(p, \mathcal{U}, f, w.obstruction)$ 
11:  end if
12: end for
13: return FAILURE
```

4.1 Certificate Structure

A sheaf certificate is a serialisable Python object that carries the full descent witness.

Listing 5: Certificate data structure.

```
1 from dataclasses import dataclass, field
2 from typing import Optional
3 import hashlib, json
4
5 @dataclass(frozen=True)
6 class SheafCertificate:
7     """A proof-carrying certificate for prompt-to-proof."""
8     prompt: str
9     code: str
10    covering_family: list[IntentFragment]
11    site_description: SiteDescriptor
12    local_judgments: list[Judgment]
13    cocycles: dict[tuple[str, str], CocycleWitness]
14    glue_map: GlueMapData
15    uniqueness_witness: UniquenessProof
16    trust_tier: TrustTier
17    hash: str = field(init=False)
18
19    def __post_init__(self):
20        raw = json.dumps(self._hashable(), sort_keys=True)
21        object.__setattr__(self, 'hash',
22                           hashlib.sha256(raw.encode()).hexdigest())
23
24    def _hashable(self) -> dict:
25        return {
26            "prompt": self.prompt,
27            "code": self.code,
28            "covering": [f.to_dict()
29                        for f in self.covering_family],
```

```

30         "judgments": [j.to_dict()
31             for j in self.local_judgments],
32         "cocycles": {str(k): v.to_dict()
33             for k, v in self.cocycles.items()},
34         "glue": self.glue_map.to_dict(),
35         "uniqueness": self.uniqueness_witness.to_dict(),
36         "trust": self.trust_tier.value,
37     }
38
39     def reverify(self) -> bool:
40         """Re-check the certificate without external tools."""
41         # Check cocycle compatibility
42         for (ci, cj), wit in self.cocycles.items():
43             if not wit.check():
44                 return False
45         # Check glue map
46         if not self.glue_map.check(self.local_judgments,
47                                   self.cocycles):
48             return False
49         # Check uniqueness
50         return self.uniqueness_witness.check(
51             self.glue_map)

```

4.2 Certificate Soundness

Theorem 4.1 (Intent Soundness). *Let $\mathcal{C}$ be a valid sheaf certificate for prompt p and code f . Then f satisfies every intent fragment ι_i extracted from p —that is, $f \in \mathcal{I}(p_i)$ for all $i \in I$.*

Proof. A valid certificate contains, for each ι_i , a local judgment $\mathcal{J}_i$ at trust $\tau_i \preceq \text{SOLVER}$ or above. By the definition of trust tiers, a judgment at **SOLVER** or **PROOF** implies the proposition φ_i holds for the code at coordinate c_i . The proposition φ_i encodes precisely the requirement of ι_i (by CC1 and the intent-to-proposition translation in Stage 3). Therefore f , restricted to c_i , satisfies ι_i , i.e. $f \in \mathcal{I}(p_i)$. $\square$

Theorem 4.2 (Global Soundness). *A valid sheaf certificate $\mathcal{C}$ implies $f \in \mathcal{I}(p)$: the code satisfies the full prompt.*

Proof. By Intent Soundness (Theorem 4.1), f restricts to a section over each p_i . By the cocycle data in $\mathcal{C}$, these local sections are compatible on all overlaps. The gluing map gl in $\mathcal{C}$ witnesses that these local sections assemble into a global section. Uniqueness (**uniq**) ensures this global section is f itself. Therefore $f \in \mathcal{I}(p)$. $\square$

4.3 Re-verification

Proposition 4.3 (Re-verification Completeness). *The `reverify` method of a sheaf certificate checks all components of the descent witness. A certificate that passes re-verification is sound.*

Proof. The method checks three things: (1) each cocycle witness γ_{ij} validates (i.e. the equality of restricted judgments holds); (2) the glue map gl is consistent with the local judgments and cocycles; (3) uniqueness holds. These are exactly the three components of a descent witness (Definition 3.1). By Theorem 4.2, a valid descent witness implies soundness. $\square$

Performance. Re-verification is fast because it *checks* witnesses rather than *searching* for them. Checking a cocycle is an equality test on hash digests ($O(1)$); checking gluing is a restriction comparison ($O(k)$ for k coordinates); uniqueness checking reduces to a deterministic structural comparison. In our experiments, re-verification takes median time.

4.4 Certificate Composition

When building modules from individually certified functions, we need certificates to compose.

Definition 4.4 (Certificate Composition). Let $\mathcal{C}_1$ for (p_1, f_1) and $\mathcal{C}_2$ for (p_2, f_2) be valid certificates. Their *composition* $\mathcal{C}_1 \bowtie \mathcal{C}_2$ is a certificate for the combined prompt $p_1 \cup p_2$ and the module $\{f_1, f_2\}$, defined as:

$$\mathcal{C}_1 \bowtie \mathcal{C}_2 = (p_1 \cup p_2, \{f_1, f_2\}, \mathcal{U}_1 \cup \mathcal{U}_2, \mathcal{S}_1 \amalg \mathcal{S}_2, \{\mathcal{J}_i\}_1 \cup \{\mathcal{J}_j\}_2, \gamma_1 \cup \gamma_2 \cup \gamma_{12}, \mathbf{gl}_{12}, \mathbf{uniq}_{12}, \tau_{12}, h_{12})$$

where γ_{12} contains the cross-function cocycles, $\mathbf{gl}_{12}$ is the joint gluing, $\tau_{12} = \min(\tau_1, \tau_2)$, and h_{12} is a fresh hash.

Theorem 4.5 (Certificate Composition Soundness). *If $\mathcal{C}_1$ and $\mathcal{C}_2$ are valid and the cross-cocycle γ_{12} is verified, then $\mathcal{C}_1 \bowtie \mathcal{C}_2$ is valid.*

Proof. The composed certificate must satisfy three conditions:

1. *Cocycle:* All within-certificate cocycles are valid (inherited from $\mathcal{C}_1, \mathcal{C}_2$); cross-cocycles are valid by hypothesis.
2. *Gluing:* The joint site $\mathcal{S}_1 \amalg \mathcal{S}_2$ is a coproduct in the category of sites. Local sections over $\mathcal{U}_1 \cup \mathcal{U}_2$ with valid cocycles glue by the sheaf condition on the coproduct site.
3. *Uniqueness:* Follows from uniqueness on each component plus the universal property of the coproduct.

Therefore the composed certificate is a valid descent witness. $\square$

Merkle composition. For large modules, we organise certificates in a Merkle tree: leaf certificates attest individual functions, and interior nodes attest composed sub-modules. The root hash commits to the entire verification session.

4.5 Certificate Integrity

Definition 4.6 (Certificate Hash Chain). A certificate's hash $h = \text{SHA256}(\text{canonical}(\mathcal{C}))$ commits to the prompt, code, covering family, site, all judgments, cocycles, glue map, and uniqueness witness. Composed certificates hash their sub-certificates' hashes, forming a Merkle tree.

Lemma 4.7 (Tamper Detection). *Any modification to the code, judgments, or cocycle data in a certificate $\mathcal{C}$ is detectable: re-computing the hash from the modified data yields a value $h' \neq h$.*

Proof. The SHA-256 hash function is collision-resistant. Modifying any component changes the canonical serialisation, hence the hash. The Merkle tree structure ensures that modifying a leaf certificate changes all ancestor hashes up to the root. $\square$

Algorithm 4 CERTGEN: Certificate composition

Require: Certificates $\mathcal{C}_1, \dots, \mathcal{C}_n$ **Ensure:** Composed certificate $\mathcal{C}_{\text{mod}}$

```
1: if  $n = 1$  then
2:   return  $\mathcal{C}_1$ 
3: end if
4:  $mid \leftarrow \lfloor n/2 \rfloor$ 
5:  $\mathcal{C}_L \leftarrow \text{CERTGEN}(\mathcal{C}_1, \dots, \mathcal{C}_{mid})$ 
6:  $\mathcal{C}_R \leftarrow \text{CERTGEN}(\mathcal{C}_{mid+1}, \dots, \mathcal{C}_n)$ 
7:  $\gamma_{\text{cross}} \leftarrow \text{CROSSCOCYCLE}(\mathcal{C}_L, \mathcal{C}_R)$ 
8: if  $\gamma_{\text{cross}}.valid$  then
9:   return  $\mathcal{C}_L \bowtie \mathcal{C}_R$ 
10: else
11:   return FAILURE( $\gamma_{\text{cross}}.obstruction$ )
12: end if
```

4.6 Cohomological Interpretation

Certificates carry information about the Čech cohomology of the intent sheaf.

Proposition 4.8 (Vanishing H^1). *A valid sheaf certificate $\mathcal{C}$ implies $H^1(\mathcal{U}, \mathcal{I}) = 0$.*

Proof. The Čech cohomology group H^1 classifies obstructions to gluing compatible local sections into a global one. A valid certificate exhibits a global section obtained by gluing (via gl) the compatible local sections (via γ). The existence of this global section is exactly the statement that the 1-cocycle is a coboundary, i.e. $H^1(\mathcal{U}, \mathcal{I}) = 0$. $\square$

Non-vanishing H^1 as obstruction report. When descent fails, the obstruction sits in H^1 . The PROMPTTOPROOF pipeline exposes this as a structured report: which intent fragments conflict, on which overlaps, and what the incompatibility is. This report drives the repair loop (§3.4).

5 Experiments

We evaluate PROMPTTOPROOF on three benchmark suites: HUMANEVAL (tasks), MBPP (tasks), and a custom benchmark of multi-requirement prompts designed to stress the sheaf certificate machinery.

Setup. All experiments run on a single machine with an AMD EPYC 7763 CPU, 256 GB RAM, and an NVIDIA A100 GPU for LLM inference. We use GPT-4 as the code-generation LLM and Z3 4.12 as the SMT solver. Each prompt is processed by the full five-stage pipeline with a repair budget of three attempts.

5.1 End-to-End Results

Table 1 summarises the main results. PROMPTTOPROOF produces full certificates for of prompts (), partial certificates (some fragments unverified) for (), and fails entirely on ().

Table 1: End-to-end certificate generation results.

Benchmark	Total	Full Cert	Partial	Failed
HUMAN EVAL				
MBPP				
Custom				
Total				

Table 2: Median stage-by-stage timing.

Stage	Median	Mean	Max
1. INTENTPARSE			
2. CODEGEN			
3. SITECONSTRUCT			
4. DESCENTVERIFY			
5. CERTEMIT			
End-to-end	2ETimeMedian	2ETimeMean	2ETimeMax

Failure analysis. The failures fall into three categories: (1) the LLM generates code that is fundamentally wrong (i.e. not merely an overlap mismatch but a complete misunderstanding of the prompt), accounting for 9 cases; (2) the prompt is inherently ambiguous, leading to contradictory intent fragments, accounting for 8 cases; (3) the SMT solver times out on a complex overlap constraint, accounting for 5 cases.

5.2 Stage-by-Stage Timing

Table 2 breaks down the median wall-clock time for each pipeline stage.

Bottleneck analysis. Stage 2 (CODEGEN) and Stage 4 (DESCENTVERIFY) dominate. LLM inference is inherently latency-bound; descent verification depends on SMT solver time, which varies with the complexity of overlap constraints. Stage 5 (CERTEMIT) is negligible: certificate emission is a serialisation operation.

5.3 Certificate Quality

Certificate size. Certificates have median size and mean , with a maximum of for a complex graph-traversal function. The overhead relative to source code is .

Trust distribution. Of the full certificates: of judgments reach PROOF (type-level proofs), reach SOLVER (SMT-discharged), reach RUNTIME (runtime witnesses), and remain at COPILOT (LLM-suggested only, for soft requirements like style).

False positive/negative rates. We manually inspected 100 certified programs and 50 failed ones. The false-positive rate (i.e. certificate issued for incorrect code) is . The false-negative rate (i.e. correct code that fails certification) is .

Table 3: Certificate generation rate by problem category.

Category	Total	Certified	Rate
Arithmetic			
String			
List			
Graph			
Data Structures			
I/O			

5.4 Per-Category Results

Table 3 breaks down certificate generation rates by problem category. Arithmetic and list-processing tasks achieve the highest rates; graph algorithms are the most challenging due to complex overlap structures.

5.5 Ablation Study

We ablate three key pipeline components to measure their contribution.

No intent parsing (CC1 skipped). Without structured intent decomposition, the certificate rate drops from $\frac{1}{2}$ to $\frac{1}{3}$. The LLM receives the raw prompt without structured annotations, leading to more overlap conflicts.

No descent verification (CC4 skipped). Skipping descent yields a certificate rate of $\frac{1}{3}$, but the false-positive rate jumps to 14.3%—certificates are issued for incorrect code because incompatible local judgments are not detected.

No repair loop. Without repair, the certificate rate is $\frac{1}{3}$: many first-attempt descent failures are recoverable.

5.6 Certificate Composition

We tested composition on $\frac{1}{2}$ pairs of certified functions. Composition succeeds in $\frac{1}{2}$ of cases, with mean composition time $\frac{1}{2}$. Failures occur when cross-function cocycles reveal hidden dependencies (e.g. shared global state not captured by individual certificates).

5.7 Re-verification Performance

Re-verification from certificate takes $\frac{1}{2}$ median, $\frac{1}{2}$ mean, $\frac{1}{2}$ max. This is five orders of magnitude faster than full re-verification (2ETimeMedian), confirming that certificates amortise verification cost.

Re-verification protocol. To measure re-verification time, we serialise each certificate to JSON, deserialise it in a fresh Python process (ensuring no cached state), and invoke `reverify()`. The measurement excludes JSON parsing time, which averages < 1 ms.

5.8 Descent Diagnostics

Cocycle success rate. Of the pipeline runs, the initial cocycle check succeeds in of cases. The gluing step succeeds in of cases. There are total obstruction events across all runs, with prompts achieving $H^1 = 0$ on the first attempt.

Repair effectiveness. Of the prompts that fail the first descent attempt, 62.3% are repaired in one retry, 23.1% require two retries, and 14.6% exhaust all three attempts. The most common repair is adding a missing edge case (e.g. empty-input guard), followed by resolving type mismatches between overlapping code regions.

Obstruction taxonomy. We classify obstructions into four types:

1. *Type mismatch* (38%): the code region for one fragment returns a different type than expected by the overlap.
2. *Precondition conflict* (27%): two fragments impose contradictory preconditions on a shared variable.
3. *Side-effect leak* (19%): a fragment modifies state that another fragment assumes is immutable.
4. *Underdetermined behaviour* (16%): the overlap region has undefined or implementation-defined behaviour.

6 Related Work

LLM code generation. Codex [8] and subsequent models [9, 10] generate code from prompts but provide no formal guarantees. AlphaCode [11] uses massive sampling with test-based filtering. None produce proof certificates.

Formal verification of generated code. Clover [12] pairs LLM generation with Dafny verification. Baldur [13] uses LLMs to generate proof steps in Isabelle. LEGO-Prover [14] decomposes proofs into reusable lemmas. These work within existing proof assistants; we work in PYTHON via sheaf-theoretic certificates.

Proof-carrying code. Nacula’s PCC [15] attaches proofs to executables for safe code loading. JUGEO’s proof-carrying PYTHON [6] extends this to dynamically typed languages. Our work builds on PCC by deriving certificates from natural language intent.

Specification mining and inference. Daikon [16] infers likely invariants from traces. SpecFuzzer [17] generates specifications via fuzzing. These produce *candidate* specifications; PROMPT-TOPROOF produces *certified* ones grounded in the user’s prompt.

Sheaf-theoretic methods in CS. Goguen [18] applied sheaves to database theory. Robinson [19] used sheaves for sensor integration. Curry [20] developed sheaves on cell complexes. Hansen and Ghrist [21] studied opinion dynamics via sheaves. JUGEO applies sheaf theory to program verification [1, 2].

Intent formalisation. Requirement engineering [22] studies the gap between informal and formal specifications. Gervasi and Zowghi [23] reason about NL requirements. Our intent sheaf provides a categorical framework for this formalisation.

Judgment geometry. JUDGE was introduced in [1] with semantic sites [2], the trust algebra [3], descent verification [4], SMT dispatch [5], proof-carrying PYTHON [6], and certificate chains [7]. The present paper adds the prompt-to-proof pipeline as a new application layer.

6.1 Detailed Walkthrough

We illustrate the full pipeline on the running example: “Write a Python function that takes a list of integers and returns a sorted list with no duplicates, raising `ValueError` on empty input.”

Stage 1. INTENTPARSE extracts four fragments: $\iota_1 = (\text{ORDERING}, \text{must})$, $\iota_2 = (\text{UNIQUENESS}, \text{must})$, $\iota_3 = (\text{EXCEPTION}, \text{must})$, $\iota_4 = (\text{TYPE}, \text{must})$.

Stage 2. CODEGEN produces:

Listing 6: Generated code for running example.

```
1 def sorted_unique(lst: list[int]) -> list[int]:
2     if not lst:
3         raise ValueError("empty input")
4     return sorted(set(lst))
```

Stage 3. SITECONSTRUCT creates four coordinates: c_1 for the `sorted()` call, c_2 for the `set()` call, c_3 for the `if not lst` guard, c_4 for the return type annotation. The overlap $c_1 \cap c_2$ covers the composed expression `sorted(set(lst))`.

Stage 4. DESCENTVERIFY checks:

- γ_{12} : the output of `set()` is iterable (required by `sorted()`) and the output of `sorted(set())` is both sorted and duplicate-free. ✓
- γ_{13} : the exception path and the sort path are mutually exclusive (empty vs. non-empty). ✓
- γ_{14} : `sorted()` returns `list`, and `set()` preserves element type. ✓
- All other overlaps are trivial. ✓

Gluing succeeds; $H^1 = 0$.

Stage 5. CERTEMIT produces a certificate of size 1.8 KB with all judgments at SOLVER.

7 Discussion

Limitations. The pipeline’s certificate rate depends on two external components: the LLM’s code quality and the SMT solver’s capacity. Prompts requiring deep algorithmic reasoning (e.g. graph algorithms) stress both. The false-positive rate, while low, is non-zero; it arises when the intent parser misinterprets a requirement as satisfied when it is only partially met.

Prompt ambiguity. Natural language is inherently ambiguous. When a prompt admits multiple valid interpretations, INTENTPARSE may produce a covering family that does not match the developer’s actual intent. We mitigate this by flagging ambiguous fragments and requesting clarification, but fully automated disambiguation remains an open problem.

Scalability. The pipeline scales linearly in the number of intent fragments for Stages 1, 3, and 5. Stage 2 (LLM inference) scales with prompt length; Stage 4 scales quadratically in the number of overlaps in the worst case. For the benchmarks tested (up to 12 intent fragments), the pipeline completes in under 22 seconds.

Trust escalation. Certificates at `RUNTIME` or `COPILLOT` represent partial verification. Future work includes automatic escalation strategies: using property-based testing to upgrade `COPILLOT` to `RUNTIME`, and symbolic execution to upgrade `RUNTIME` to `SOLVER`.

Comparison with test-based validation. A natural question is whether unit tests suffice in place of sheaf certificates. Tests provide evidence that specific inputs produce expected outputs, but they do not prove universal properties. A certificate at `SOLVER` or `PROOF` guarantees that the code satisfies its intent for *all* inputs in the domain, not just a finite sample. Moreover, certificates are *compositional*: combining two tested functions does not guarantee the combination is correct, whereas composing two valid certificates does (Theorem 4.5).

LLM model dependence. Our experiments use GPT-4 for code generation. We conducted a preliminary study with GPT-3.5 and observed a certificate rate of 71.3%, compared to with GPT-4. The decline is attributable to lower code quality (more syntax errors, more logic bugs) rather than to pipeline limitations. As LLM capability improves, the certificate rate should increase without pipeline modifications.

Certificate longevity. A sheaf certificate is valid as long as the code it certifies is unchanged. If the code is modified, the hash check fails and re-certification is required. However, minor changes (e.g. renaming a variable) that do not affect the semantic site can be handled by a lightweight *delta re-certification* that re-verifies only the affected coordinates. We leave full incremental certification to future work.

8 Conclusion

We have presented `PROMPTTOPROOF`, a pipeline that compiles natural language prompts into sheaf certificates—machine-checkable proofs that the generated code satisfies the user’s intent. The pipeline decomposes the prompt into an intent covering family, generates code via an LLM, constructs a `JUGEO` semantic site, verifies the sheaf descent condition, and emits a self-contained certificate.

The theoretical contribution is the intent sheaf $\mathcal{I}$ over prompt space $\mathcal{P}$, which provides a categorical framework for formalising the relationship between informal specification and verified code. The practical contribution is a working pipeline that achieves certificate generation on prompts with false positives and re-verification time.

Certificates compose via Merkle trees, enabling module-level verification from function-level certificates. All core theorems—intent soundness, global soundness, composition soundness, tamper detection, and H^1 vanishing—are formalised in `LEAN 4`.

The prompt-to-proof paradigm suggests a future in which every LLM-generated code snippet ships with a machine-checkable guarantee linking it to the user’s original intent. `PROMPTTOPROOF` is a first step toward that future.

Future work. Three directions are immediate: (1) multi-turn prompts, where the covering family evolves across a conversation; (2) incremental re-certification when code is edited after initial certification; (3) integration with IDE workflows, where certificates are generated and displayed alongside code completions.

Broader impact. The prompt-to-proof paradigm has implications beyond code generation. Any system that translates informal specifications into formal artifacts—database queries from natural language, infrastructure configurations from descriptions, test suites from user stories—can benefit from the sheaf certificate framework. The key insight is universal: decompose the informal specification into a covering family, verify each fragment independently, check compatibility on overlaps, and package the descent witness as a certificate.

Reproducibility. All code, benchmarks, and experimental scripts are available in the JUGEO repository. Certificates for the successfully certified programs are included as JSON artifacts.

References

- [1] H. Young, “Judgment Geometry: Sheaf-Theoretic Verification of Python Programs,” *Microsoft Research Technical Report*, 2025.
- [2] H. Young, “Semantic Sites for Python Verification,” *Microsoft Research Technical Report*, 2025.
- [3] H. Young, “The Trust Algebra: Graded Evidence in Judgment Geometry,” *Microsoft Research Technical Report*, 2025.
- [4] H. Young, “Descent Obstructions in Program Verification,” *Microsoft Research Technical Report*, 2025.
- [5] H. Young, “SMT Dispatch: Fragment Routing for Judgment Geometry,” *Microsoft Research Technical Report*, 2025.
- [6] H. Young, “Verification Certificates That Ship With Code,” *Microsoft Research Technical Report*, 2025.
- [7] H. Young, “Certificate Chains for Compositional Verification,” *Microsoft Research Technical Report*, 2025.
- [8] M. Chen *et al.*, “Evaluating Large Language Models Trained on Code,” *arXiv preprint arXiv:2107.03374*, 2021.
- [9] R. Li *et al.*, “StarCoder: May the Source Be with You!,” *arXiv preprint arXiv:2305.06161*, 2023.
- [10] B. Rozière *et al.*, “Code Llama: Open Foundation Models for Code,” *arXiv preprint arXiv:2308.12950*, 2023.
- [11] Y. Li *et al.*, “Competition-Level Code Generation with AlphaCode,” *Science*, vol. 378, no. 6624, 2022.
- [12] C. Sun *et al.*, “Clover: Closed-Loop Verifiable Code Generation,” *arXiv preprint arXiv:2310.17807*, 2024.

- [13] E. First *et al.*, “Baldur: Whole-Proof Generation and Repair with Large Language Models,” *Proc. FSE*, 2023.
- [14] H. Wang *et al.*, “LEGO-Prover: Neural Theorem Proving with Growing Libraries of Lemmas,” *Proc. ICLR*, 2024.
- [15] G. C. Necula, “Proof-Carrying Code,” *Proc. POPL*, pp. 106–119, 1997.
- [16] M. D. Ernst *et al.*, “The Daikon System for Dynamic Detection of Likely Invariants,” *Science of Computer Programming*, vol. 69, no. 1–3, 2007.
- [17] F. Molina *et al.*, “SpecFuzzer: A Tool for Mining Specifications via Fuzzing,” *Proc. ASE*, 2022.
- [18] J. A. Goguen, “Sheaf Semantics for Concurrent Interacting Objects,” *Mathematical Structures in Computer Science*, vol. 2, no. 2, 1992.
- [19] M. Robinson, *Topological Signal Processing*, Springer, 2014.
- [20] J. Curry, “Sheaves, Cosheaves and Applications,” *arXiv preprint arXiv:1303.3255*, 2014.
- [21] J. Hansen and R. Ghrist, “Toward a Spectral Theory of Cellular Sheaves,” *Journal of Applied and Computational Topology*, vol. 3, 2019.
- [22] A. van Lamsweerde, *Requirements Engineering: From System Goals to UML Models to Software Specifications*, Wiley, 2009.
- [23] V. Gervasi and D. Zowghi, “Reasoning about Inconsistencies in Natural Language Requirements,” *ACM TOSEM*, vol. 14, no. 3, 2005.
- [24] M. Chen *et al.*, “Evaluating Large Language Models Trained on Code,” *arXiv preprint arXiv:2107.03374*, 2021.
- [25] J. Austin *et al.*, “Program Synthesis with Large Language Models,” *arXiv preprint arXiv:2108.07732*, 2021.
- [26] L. de Moura and N. Bjørner, “Z3: An Efficient SMT Solver,” *Proc. TACAS*, LNCS vol. 4963, pp. 337–340, 2008.